

### Problem Statement:

- **Classification problem**
  - Collect dataset from here  
<https://archive.ics.uci.edu/ml/datasets/census+income>
  - Here we have missing values also.
  - So here perform EDA, Data wrangling, Data Pre processing
  - Now make a Classification model to find how many people are >50k and how many are <=50k.
  - Here create Logistic Regression, SVM.
- **Steps to be followed**
  - Data ingestion.
  - EDA (end to end).
  - Preprocessing of the data.
  - Use pickle to store the scaling of the data for later use.
  - Store the final processed data inside MongoDB.
  - Again load the data from MongoDB.
  - Model building.
  - Use GridSearchCV for hyper parameter tuning.
  - Evaluation.
    - Confusion Matrix, ROC and AUC for classification model.

### Attribute Information:

1. **age:** continuous.
2. **workclass:** Private, Self-emp-not-inc, Self-emp-inc, Federal-gov, Local-gov, State-gov, Without-pay, Never-worked.
3. **fnlwgt:** continuous.
4. **education:** Bachelors, Some-college, 11th, HS-grad, Prof-school, Assoc-acdm, Assoc-voc, 9th, 7th-8th, 12th, Masters, 1st-4th, 10th, Doctorate, 5th-6th, Preschool.
5. **education-num:** continuous.
6. **marital-status:** Married-civ-spouse, Divorced, Never-married, Separated, Widowed, Married-spouse-absent, Married-AF-spouse.
7. **occupation:** Tech-support, Craft-repair, Other-service, Sales, Exec-managerial, Prof-specialty, Handlers-cleaners, Machine-op-inspct, Adm-clerical, Farming-fishing, Transport-moving, Priv-house-serv, Protective-serv, Armed-Forces.
8. **relationship:** Wife, Own-child, Husband, Not-in-family, Other-relative, Unmarried.
9. **race:** White, Asian-Pac-Islander, Amer-Indian-Eskimo, Other, Black.
10. **sex:** Female, Male.
11. **capital-gain:** continuous.
12. **capital-loss:** continuous.
13. **hours-per-week:** continuous.

14. **native-country:** United-States, Cambodia, England, Puerto-Rico, Canada, Germany, Outlying-US(Guam-USVI-etc), India, Japan, Greece, South, China, Cuba, Iran, Honduras, Philippines, Italy, Poland, Jamaica, Vietnam, Mexico, Portugal, Ireland, France, Dominican-Republic, Laos, Ecuador, Taiwan, Haiti, Columbia, Hungary, Guatemala, Nicaragua, Scotland, Thailand, Yugoslavia, El-Salvador, Trinidad&Tobago, Peru, Hong, Holand-Netherlands.

## 1. Data Ingestion:

### 1.1 Import modules and data to create dataframe

*# Importing the required libraries*

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
import pymongo
```

```
sns.set()
%matplotlib inline
warnings.filterwarnings('ignore')
```

*# Creating final csv file*

```
df = pd.read_csv('dataset/adult.csv', header=None)
df.rename(columns={0:'age', 1:'workclass', 2:'fnlwgt', 3:'education',
4:'education-num', 5:'marital-status',
6:'occupation', 7:'relationship', 8:'race',
9:'sex', 10:'capital-gain', 11:'capital-loss',
12:'hours-per-week', 13:'native-country',
14:'income'}, inplace=True)
df.to_csv('dataset/adult_final.csv', index=False)
```

### 1.2 Creating dataframe

```
data = pd.read_csv('dataset/adult_final.csv')
data.head()
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	marital-status	occupation	relationship	race
sex \				
0	Never-married	Adm-clerical	Not-in-family	White
Male				
1	Married-civ-spouse	Exec-managerial	Husband	White

```

Male
2          Divorced   Handlers-cleaners   Not-in-family   White
Male
3   Married-civ-spouse   Handlers-cleaners           Husband   Black
Male
4   Married-civ-spouse           Prof-specialty           Wife   Black
Female

```

```

      capital-gain  capital-loss  hours-per-week  native-country  income
0          2174           0           40   United-States  <=50K
1           0           0           13   United-States  <=50K
2           0           0           40   United-States  <=50K
3           0           0           40   United-States  <=50K
4           0           0           40           Cuba  <=50K

```

```
data.shape
```

```
(32561, 15)
```

```
# Checking the data types
```

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 32561 entries, 0 to 32560
```

```
Data columns (total 15 columns):
```

#	Column	Non-Null Count	Dtype
0	age	32561 non-null	int64
1	workclass	32561 non-null	object
2	fnlwgt	32561 non-null	int64
3	education	32561 non-null	object
4	education-num	32561 non-null	int64
5	marital-status	32561 non-null	object
6	occupation	32561 non-null	object
7	relationship	32561 non-null	object
8	race	32561 non-null	object
9	sex	32561 non-null	object
10	capital-gain	32561 non-null	int64
11	capital-loss	32561 non-null	int64
12	hours-per-week	32561 non-null	int64
13	native-country	32561 non-null	object
14	income	32561 non-null	object

```
dtypes: int64(6), object(9)
```

```
memory usage: 3.7+ MB
```

## Observations:

- There are 32561 rows with 15 columns.

## 2. Data Cleaning

*# Name of the columns*

```
data.columns
```

```
Index(['age', 'workclass', 'fnlwgt', 'education', 'education-num',  
      'marital-status', 'occupation', 'relationship', 'race', 'sex',  
      'capital-gain', 'capital-loss', 'hours-per-week', 'native-  
country',  
      'income'],  
      dtype='object')
```

### 2.1 Checking all the unique values in each columns

```
for column in data.columns:  
    print(f"The unique values in column {column}:")  
    print(data[column].unique())  
    print(f"\nThe number of unique values in {column} is:  
{len(data[column].unique())}")  
    print("-----\n")
```

The unique values in column age:

```
[39 50 38 53 28 37 49 52 31 42 30 23 32 40 34 25 43 54 35 59 56 19 20  
45  
22 48 21 24 57 44 41 29 18 47 46 36 79 27 67 33 76 17 55 61 70 64 71  
68  
66 51 58 26 60 90 75 65 77 62 63 80 72 74 69 73 81 78 88 82 83 84 85  
86  
87]
```

The number of unique values in age is: 73

-----

The unique values in column workclass:

```
[' State-gov' ' Self-emp-not-inc' ' Private' ' Federal-gov' ' Local-  
gov'  
 ' ?' ' Self-emp-inc' ' Without-pay' ' Never-worked']
```

The number of unique values in workclass is: 9

-----

The unique values in column fnlwgt:

```
[ 77516  83311 215646 ... 34066  84661 257302]
```

The number of unique values in fnlwgt is: 21648

-----

The unique values in column education:

```
[' Bachelors' ' HS-grad' ' 11th' ' Masters' ' 9th' ' Some-college'  
 ' Assoc-acdm' ' Assoc-voc' ' 7th-8th' ' Doctorate' ' Prof-school'  
 ' 5th-6th' ' 10th' ' 1st-4th' ' Preschool' ' 12th']
```

The number of unique values in education is: 16

-----

The unique values in column education-num:

```
[13  9  7 14  5 10 12 11  4 16 15  3  6  2  1  8]
```

The number of unique values in education-num is: 16

-----

The unique values in column marital-status:

```
[' Never-married' ' Married-civ-spouse' ' Divorced'  
 ' Married-spouse-absent' ' Separated' ' Married-AF-spouse' '  
Widowed']
```

The number of unique values in marital-status is: 7

-----

The unique values in column occupation:

```
[' Adm-clerical' ' Exec-managerial' ' Handlers-cleaners' ' Prof-  
specialty'  
 ' Other-service' ' Sales' ' Craft-repair' ' Transport-moving'  
 ' Farming-fishing' ' Machine-op-inspct' ' Tech-support' ' ?'  
 ' Protective-serv' ' Armed-Forces' ' Priv-house-serv']
```

The number of unique values in occupation is: 15

-----

The unique values in column relationship:

```
[' Not-in-family' ' Husband' ' Wife' ' Own-child' ' Unmarried'  
 ' Other-relative']
```

The number of unique values in relationship is: 6

-----

The unique values in column race:

```
[' White' ' Black' ' Asian-Pac-Islander' ' Amer-Indian-Eskimo' '  
Other']
```

The number of unique values in race is: 5

-----

The unique values in column sex:

```
[' Male' ' Female']
```

The number of unique values in sex is: 2

-----

The unique values in column capital-gain:

```
[ 2174      0 14084  5178  5013  2407 14344 15024  7688 34095  4064
4386
  7298  1409  3674  1055  3464  2050  2176   594 20051  6849  4101
1111
  8614  3411  2597 25236  4650  9386  2463  3103 10605  2964  3325
2580
  3471  4865 99999  6514  1471  2329  2105  2885 25124 10520  2202
2961
 27828  6767  2228  1506 13550  2635  5556  4787  3781  3137  3818
3942
   914   401  2829  2977  4934  2062  2354  5455 15020  1424  3273
22040
  4416  3908 10566   991  4931  1086  7430  6497   114  7896  2346
3418
  3432  2907  1151  2414  2290 15831 41310  4508  2538  3456  6418
1848
  3887  5721  9562  1455  2036  1831 11678  2936  2993  7443  6360
1797
  1173  4687  6723  2009  6097  2653  1639 18481  7978  2387  5060]
```

The number of unique values in capital-gain is: 119

-----

The unique values in column capital-loss:

```
[    0 2042 1408 1902 1573 1887 1719 1762 1564 2179 1816 1980 1977 1876
1340 2206 1741 1485 2339 2415 1380 1721 2051 2377 1669 2352 1672  653
2392 1504 2001 1590 1651 1628 1848 1740 2002 1579 2258 1602  419 2547
2174 2205 1726 2444 1138 2238  625  213 1539  880 1668 1092 1594 3004
2231 1844  810 2824 2559 2057 1974  974 2149 1825 1735 1258 2129 2603
2282  323 4356 2246 1617 1648 2489 3770 1755 3683 2267 2080 2457  155
3900 2201 1944 2467 2163 2754 2472 1411]
```

The number of unique values in capital-loss is: 92

-----

The unique values in column hours-per-week:

```
[40 13 16 45 50 80 30 35 60 20 52 44 15 25 38 43 55 48 58 32 70  2 22
56
 41 28 36 24 46 42 12 65  1 10 34 75 98 33 54  8  6 64 19 18 72  5  9
47
 37 21 26 14  4 59  7 99 53 39 62 57 78 90 66 11 49 84  3 17 68 27 85
31
 51 77 63 23 87 88 73 89 97 94 29 96 67 82 86 91 81 76 92 61 74 95]
```

The number of unique values in hours-per-week is: 94

-----

The unique values in column native-country:

```
[' United-States' ' Cuba' ' Jamaica' ' India' ' ?' ' Mexico' ' South'
 ' Puerto-Rico' ' Honduras' ' England' ' Canada' ' Germany' ' Iran'
 ' Philippines' ' Italy' ' Poland' ' Columbia' ' Cambodia' ' Thailand'
 ' Ecuador' ' Laos' ' Taiwan' ' Haiti' ' Portugal' ' Dominican-
 Republic'
 ' El-Salvador' ' France' ' Guatemala' ' China' ' Japan' ' Yugoslavia'
 ' Peru' ' Outlying-US(Guam-USVI-etc)' ' Scotland' ' Trinidad&Tobago'
 ' Greece' ' Nicaragua' ' Vietnam' ' Hong' ' Ireland' ' Hungary'
 ' Holand-Netherlands']
```

The number of unique values in native-country is: 42

-----

The unique values in column income:

```
[' <=50K' ' >50K']
```

The number of unique values in income is: 2

-----

### Checking number of occurrence of each unique values

```
for column in data.columns:
    print(f"The number of occurrence of each values in column {column}
is:")
    print(data[column].value_counts())
    print("-----\n\n")
```

The number of occurrence of each values in column age is:

```
36      898
31      888
34      886
23      877
35      876
...
83         6
88         3
85         3
86         1
87         1
```

Name: age, Length: 73, dtype: int64

-----

The number of occurrence of each values in column workclass is:

```
Private      22696
```

Self-emp-not-inc	2541
Local-gov	2093
?	1836
State-gov	1298
Self-emp-inc	1116
Federal-gov	960
Without-pay	14
Never-worked	7

Name: workclass, dtype: int64

-----

The number of occurrence of each values in column fnlwgt is:

164190	13
203488	13
123011	13
148995	12
121124	12
..	
232784	1
325573	1
140176	1
318264	1
257302	1

Name: fnlwgt, Length: 21648, dtype: int64

-----

The number of occurrence of each values in column education is:

HS-grad	10501
Some-college	7291
Bachelors	5355
Masters	1723
Assoc-voc	1382
11th	1175
Assoc-acdm	1067
10th	933
7th-8th	646
Prof-school	576
9th	514
12th	433
Doctorate	413
5th-6th	333
1st-4th	168
Preschool	51

Name: education, dtype: int64

-----

The number of occurrence of each values in column education-num is:



9	10501
10	7291
13	5355
14	1723
11	1382
7	1175
12	1067
6	933
4	646
15	576
5	514
8	433
16	413
3	333
2	168
1	51

Name: education-num, dtype: int64

-----

The number of occurrence of each values in column marital-status is:

Married-civ-spouse	14976
Never-married	10683
Divorced	4443
Separated	1025
Widowed	993
Married-spouse-absent	418
Married-AF-spouse	23

Name: marital-status, dtype: int64

-----

The number of occurrence of each values in column occupation is:

Prof-specialty	4140
Craft-repair	4099
Exec-managerial	4066
Adm-clerical	3770
Sales	3650
Other-service	3295
Machine-op-inspct	2002
?	1843
Transport-moving	1597
Handlers-cleaners	1370
Farming-fishing	994
Tech-support	928
Protective-serv	649
Priv-house-serv	149
Armed-Forces	9

Name: occupation, dtype: int64

-----

The number of occurrence of each values in column relationship is:

Husband	13193
Not-in-family	8305
Own-child	5068
Unmarried	3446
Wife	1568
Other-relative	981

Name: relationship, dtype: int64

The number of occurrence of each values in column race is:

White	27816
Black	3124
Asian-Pac-Islander	1039
Amer-Indian-Eskimo	311
Other	271

Name: race, dtype: int64

The number of occurrence of each values in column sex is:

Male	21790
Female	10771

Name: sex, dtype: int64

The number of occurrence of each values in column capital-gain is:

0	29849
15024	347
7688	284
7298	246
99999	159

...	
1111	1
2538	1
22040	1
4931	1
5060	1

Name: capital-gain, Length: 119, dtype: int64

The number of occurrence of each values in column capital-loss is:

0	31042
1902	202

```

1977      168
1887      159
1848       51
...
2080       1
1539       1
1844       1
2489       1
1411       1
Name: capital-loss, Length: 92, dtype: int64
-----

```

The number of occurrence of each values in column hours-per-week is:

```

40      15217
50      2819
45      1824
60      1475
35      1297
...
82       1
92       1
87       1
74       1
94       1
Name: hours-per-week, Length: 94, dtype: int64
-----

```

The number of occurrence of each values in column native-country is:

```

United-States      29170
Mexico              643
?                  583
Philippines        198
Germany            137
Canada             121
Puerto-Rico       114
El-Salvador        106
India              100
Cuba                95
England            90
Jamaica            81
South              80
China              75
Italy              73
Dominican-Republic 70
Vietnam            67
Guatemala          64
Japan              62
Poland             60

```

Columbia	59
Taiwan	51
Haiti	44
Iran	43
Portugal	37
Nicaragua	34
Peru	31
France	29
Greece	29
Ecuador	28
Ireland	24
Hong	20
Cambodia	19
Trinidad&Tobago	19
Laos	18
Thailand	18
Yugoslavia	16
Outlying-US(Guam-USVI-etc)	14
Honduras	13
Hungary	13
Scotland	12
Holand-Netherlands	1

Name: native-country, dtype: int64

-----

The number of occurrence of each values in column income is:

<=50K	24720
>50K	7841

Name: income, dtype: int64

-----

### Observations:

- We have special character ? in columns workclass, occupation, native-country.
- Also the column fnlwgt has more than 1000 unique values. So we need to check for special characters in that as well.

*# Searching for special character in 'fnlwgt' column*

```
data.loc[data['fnlwgt'] == ' ?', :]
```

Empty DataFrame

Columns: [age, workclass, fnlwgt, education, education-num, marital-status, occupation, relationship, race, sex, capital-gain, capital-loss, hours-per-week, native-country, income]

Index: []

## Observation:

- So there is no special character ? in the column fnlwgt.

## Replacing the special character ? with most appeared value in that column

```
data['workclass'] = data['workclass'].str.replace('?', 'Private')
data['occupation'] = data['occupation'].str.replace('?', 'Prof-
specialty')
data['native-country'] = data['native-
country'].str.replace('?', 'United-States')
data.head()
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	sex	marital-status	occupation	relationship	race
0	Male	Never-married	Adm-clerical	Not-in-family	White
1	Male	Married-civ-spouse	Exec-managerial	Husband	White
2	Male	Divorced	Handlers-cleaners	Not-in-family	White
3	Male	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	Female	Married-civ-spouse	Prof-specialty	Wife	Black

	capital-gain	capital-loss	hours-per-week	native-country	income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

*# Verifying the result*

```
data.loc[data['workclass'] == '?', :]
```

Empty DataFrame

Columns: [age, workclass, fnlwgt, education, education-num, marital-

```
status, occupation, relationship, race, sex, capital-gain, capital-  
loss, hours-per-week, native-country, income]  
Index: []
```

```
data.loc[data['occupation'] == '?', :]
```

```
Empty DataFrame
```

```
Columns: [age, workclass, fnlwgt, education, education-num, marital-  
status, occupation, relationship, race, sex, capital-gain, capital-  
loss, hours-per-week, native-country, income]  
Index: []
```

```
data.loc[data['native-country'] == '?', :]
```

```
Empty DataFrame
```

```
Columns: [age, workclass, fnlwgt, education, education-num, marital-  
status, occupation, relationship, race, sex, capital-gain, capital-  
loss, hours-per-week, native-country, income]  
Index: []
```

*Checking for null and duplicate values*

*# checking for null values*

```
data.isnull().sum()
```

```
age                0  
workclass          0  
fnlwgt             0  
education          0  
education-num      0  
marital-status     0  
occupation         0  
relationship       0  
race               0  
sex                0  
capital-gain       0  
capital-loss       0  
hours-per-week     0  
native-country     0  
income             0  
dtype: int64
```

*# checking for duplicate values*

```
df[data.duplicated()]
```

	age	workclass	fnlwgt	education	education-num	\
4881	25	Private	308144	Bachelors	13	
5104	90	Private	52386	Some-college	10	
9171	21	Private	250051	Some-college	10	
11631	20	Private	107658	Some-college	10	
13084	25	Private	195994	1st-4th	2	

15059	21	Private	243368	Preschool	1
17040	46	Private	173243	HS-grad	9
18555	30	Private	144593	HS-grad	9
18698	19	Private	97261	HS-grad	9
21318	19	Private	138153	Some-college	10
21490	19	Private	146679	Some-college	10
21875	49	Private	31267	7th-8th	4
22300	25	Private	195994	1st-4th	2
22367	44	Private	367749	Bachelors	13
22494	49	Self-emp-not-inc	43479	Some-college	10
25872	23	Private	240137	5th-6th	3
26313	28	Private	274679	Masters	14
28230	27	Private	255582	HS-grad	9
28522	42	Private	204235	Some-college	10
28846	39	Private	30916	HS-grad	9
29157	38	Private	207202	HS-grad	9
30845	46	Private	133616	Some-college	10
31993	19	Private	251579	Some-college	10
32404	35	Private	379959	HS-grad	9

	marital-status	occupation	relationship	\
4881	Never-married	Craft-repair	Not-in-family	
5104	Never-married	Other-service	Not-in-family	
9171	Never-married	Prof-specialty	Own-child	
11631	Never-married	Tech-support	Not-in-family	
13084	Never-married	Priv-house-serv	Not-in-family	
15059	Never-married	Farming-fishing	Not-in-family	
17040	Married-civ-spouse	Craft-repair	Husband	
18555	Never-married	Other-service	Not-in-family	
18698	Never-married	Farming-fishing	Not-in-family	
21318	Never-married	Adm-clerical	Own-child	
21490	Never-married	Exec-managerial	Own-child	
21875	Married-civ-spouse	Craft-repair	Husband	
22300	Never-married	Priv-house-serv	Not-in-family	
22367	Never-married	Prof-specialty	Not-in-family	
22494	Married-civ-spouse	Craft-repair	Husband	
25872	Never-married	Handlers-cleaners	Not-in-family	
26313	Never-married	Prof-specialty	Not-in-family	
28230	Never-married	Machine-op-inspct	Not-in-family	
28522	Married-civ-spouse	Prof-specialty	Husband	
28846	Married-civ-spouse	Craft-repair	Husband	
29157	Married-civ-spouse	Machine-op-inspct	Husband	
30845	Divorced	Adm-clerical	Unmarried	
31993	Never-married	Other-service	Own-child	
32404	Divorced	Other-service	Not-in-family	

	race	sex	capital-gain	capital-loss	\
4881	White	Male	0	0	
5104	Asian-Pac-Islander	Male	0	0	
9171	White	Female	0	0	

11631	White	Female	0	0
13084	White	Female	0	0
15059	White	Male	0	0
17040	White	Male	0	0
18555	Black	Male	0	0
18698	White	Male	0	0
21318	White	Female	0	0
21490	Black	Male	0	0
21875	White	Male	0	0
22300	White	Female	0	0
22367	White	Female	0	0
22494	White	Male	0	0
25872	White	Male	0	0
26313	White	Male	0	0
28230	White	Female	0	0
28522	White	Male	0	0
28846	White	Male	0	0
29157	White	Male	0	0
30845	White	Female	0	0
31993	White	Male	0	0
32404	White	Female	0	0

	hours-per-week	native-country	income
4881	40	Mexico	<=50K
5104	35	United-States	<=50K
9171	10	United-States	<=50K
11631	10	United-States	<=50K
13084	40	Guatemala	<=50K
15059	50	Mexico	<=50K
17040	40	United-States	<=50K
18555	40	?	<=50K
18698	40	United-States	<=50K
21318	10	United-States	<=50K
21490	30	United-States	<=50K
21875	40	United-States	<=50K
22300	40	Guatemala	<=50K
22367	45	Mexico	<=50K
22494	40	United-States	<=50K
25872	55	Mexico	<=50K
26313	50	United-States	<=50K
28230	40	United-States	<=50K
28522	40	United-States	>50K
28846	40	United-States	<=50K
29157	48	United-States	>50K
30845	40	United-States	<=50K
31993	14	United-States	<=50K
32404	40	United-States	<=50K



*# Number of duplicated values*

```
df[data.duplicated()].count()
```

```
age                24
workclass          24
fnlwgt             24
education          24
education-num      24
marital-status     24
occupation         24
relationship       24
race              24
sex               24
capital-gain       24
capital-loss       24
hours-per-week     24
native-country     24
income            24
dtype: int64
```

*# Dropping the duplicated values as their numbers are minimal in respect of the dataset*

```
data.drop_duplicates(inplace=True)
```

```
data.shape
```

```
(32537, 15)
```

### Observations

- So now we have 32537 rows with no null or duplicated values.

### Let's save this clean dataset for future use

```
try:
    data.to_csv("dataset/adult_cleaned.csv", index=False)
except Exception as err:
    print("Error is: ", err)
else:
    print("Clean csv file created successfully.")
```

Clean csv file created successfully.

### 3. EDA:

#### Using the clean data

```
df = pd.read_csv('dataset/adult_cleaned.csv')
df.head()
```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	sex	marital-status	occupation	relationship	race
0	Male	Never-married	Adm-clerical	Not-in-family	White
1	Male	Married-civ-spouse	Exec-managerial	Husband	White
2	Male	Divorced	Handlers-cleaners	Not-in-family	White
3	Male	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	Female	Married-civ-spouse	Prof-specialty	Wife	Black

	capital-gain	capital-loss	hours-per-week	native-country	income
0	2174	0	40	United-States	<=50K
1	0	0	13	United-States	<=50K
2	0	0	40	United-States	<=50K
3	0	0	40	United-States	<=50K
4	0	0	40	Cuba	<=50K

### 3.1 Differentiating numerical and categorical columns

```
numerical_features = [feature for feature in df.columns if
df[feature].dtypes != 'O']
categorical_features = [feature for feature in df.columns if
df[feature].dtypes == 'O']
```

```
print(f"The number of Numerical features are:
{len(numerical_features)}, and the column names are:\n
{n{numerical_features}}")
print(f"\nThe number of Categorical features are:
{len(categorical_features)}, and the column names are:\n
{n{categorical_features}}")
```

The number of Numerical features are: 6, and the column names are:  
['age', 'fnlwgt', 'education-num', 'capital-gain', 'capital-loss',  
'hours-per-week']

The number of Categorical features are: 9, and the column names are:  
['workclass', 'education', 'marital-status', 'occupation',  
'relationship', 'race', 'sex', 'native-country', 'income']

*# proportion of count data on categorical columns*

```
for col in categorical_features:  
    print(df[col].value_counts(normalize=True) * 100)  
    print('-----')
```

Private	75.326551
Self-emp-not-inc	7.806497
Local-gov	6.432677
State-gov	3.989304
Self-emp-inc	3.429941
Federal-gov	2.950487
Without-pay	0.043028
Never-worked	0.021514

Name: workclass, dtype: float64

HS-grad	32.252513
Some-college	22.380674
Bachelors	16.452039
Masters	5.292436
Assoc-voc	4.247472
11th	3.611273
Assoc-acdm	3.279344
10th	2.867505
7th-8th	1.982359
Prof-school	1.770292
9th	1.579740
12th	1.330793
Doctorate	1.269324
5th-6th	1.020377
1st-4th	0.510188
Preschool	0.153671

Name: education, dtype: float64

Married-civ-spouse	46.009159
Never-married	32.784215
Divorced	13.649076
Separated	3.150260
Widowed	3.051910
Married-spouse-absent	1.284691
Married-AF-spouse	0.070689

Name: marital-status, dtype: float64

Prof-specialty	18.376003
Craft-repair	12.582598
Exec-managerial	12.493469

Adm-clerical	11.580662
Sales	11.217998
Other-service	10.114639
Machine-op-inspct	6.146848
Transport-moving	4.908258
Handlers-cleaners	4.207518
Farming-fishing	3.048837
Tech-support	2.849064
Protective-serv	1.994652
Priv-house-serv	0.451793
Armed-Forces	0.027661

Name: occupation, dtype: float64

-----

Husband	40.529244
Not-in-family	25.484833
Own-child	15.563820
Unmarried	10.587946
Wife	4.819129
Other-relative	3.015029

Name: relationship, dtype: float64

-----

White	85.425823
Black	9.595230
Asian-Pac-Islander	3.190214
Amer-Indian-Eskimo	0.955835
Other	0.832898

Name: race, dtype: float64

-----

Male	66.92381
Female	33.07619

Name: sex, dtype: float64

-----

United-States	91.388266
Mexico	1.963918
Philippines	0.608538
Germany	0.421059
Canada	0.371884
Puerto-Rico	0.350370
El-Salvador	0.325783
India	0.307342
Cuba	0.291975
England	0.276608
Jamaica	0.248947
South	0.245874
China	0.230507
Italy	0.224360
Dominican-Republic	0.215140
Vietnam	0.205919
Japan	0.190552
Guatemala	0.190552

Poland	0.184405
Columbia	0.181332
Taiwan	0.156745
Haiti	0.135231
Iran	0.132157
Portugal	0.113717
Nicaragua	0.104496
Peru	0.095276
France	0.089129
Greece	0.089129
Ecuador	0.086056
Ireland	0.073762
Hong	0.061468
Cambodia	0.058395
Trinidad&Tobago	0.058395
Laos	0.055322
Thailand	0.055322
Yugoslavia	0.049175
Outlying-US(Guam-USVI-etc)	0.043028
Honduras	0.039955
Hungary	0.039955
Scotland	0.036881
Holand-Netherlands	0.003073

Name: native-country, dtype: float64

-----

<=50K      75.907428

>50K        24.092572

Name: income, dtype: float64

-----

### 3.2 Statistical Analysis of the data

# summary of the dataset

df.describe().T

	count	mean	std	min
25% \				
age	32537.0	38.585549	13.637984	17.0
28.0				
fnlwgt	32537.0	189780.848511	105556.471009	12285.0
117827.0				
education-num	32537.0	10.081815	2.571633	1.0
9.0				
capital-gain	32537.0	1078.443741	7387.957424	0.0
0.0				
capital-loss	32537.0	87.368227	403.101833	0.0
0.0				
hours-per-week	32537.0	40.440329	12.346889	1.0
40.0				

	50%	75%	max
age	37.0	48.0	90.0
fnlwgt	178356.0	236993.0	1484705.0
education-num	10.0	12.0	16.0
capital-gain	0.0	0.0	99999.0
capital-loss	0.0	0.0	4356.0
hours-per-week	40.0	45.0	99.0

### Observations:

- There are outliers in all the numerical columns except education-num.

## 3.3 Graphical Analysis of the data

### 3.3.1 Univariate Analysis

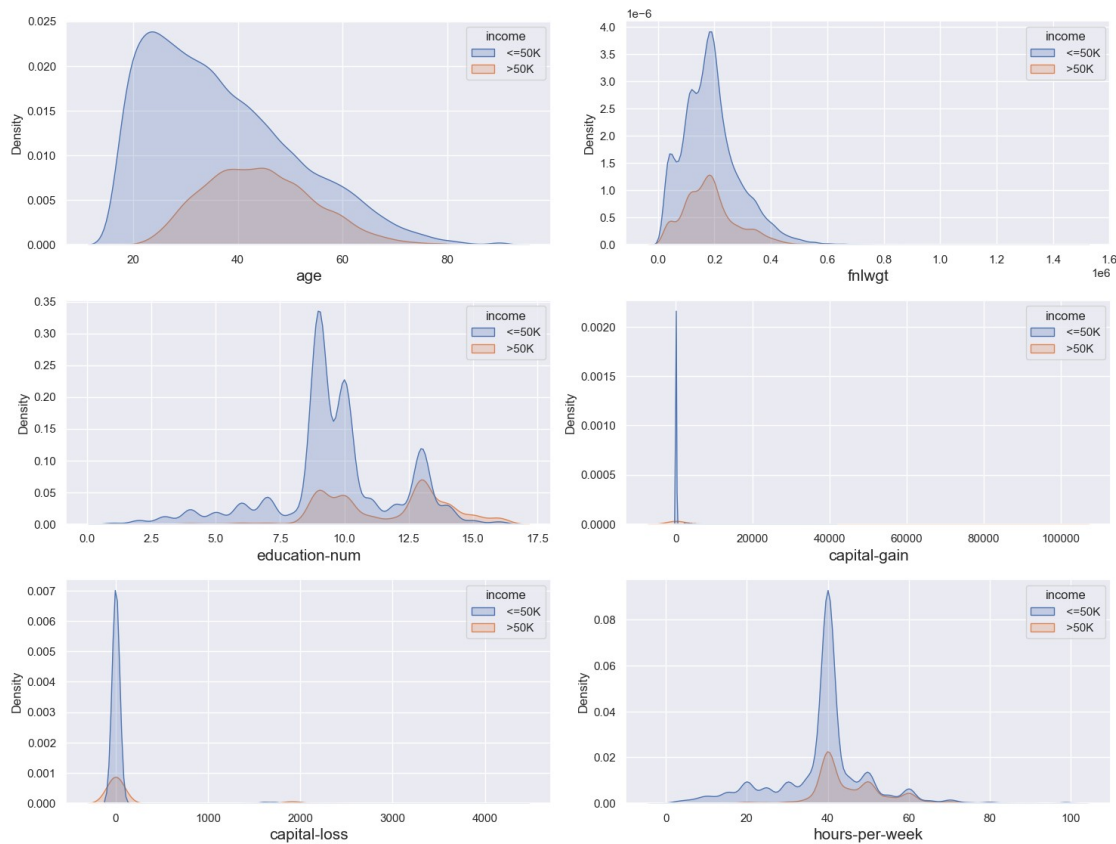
#### Numerical Features

*# Kernal Density plots*

```
plt.figure(figsize=(15, 15))
plt.suptitle('Univariate Analysis of Numerical Features', fontsize=20,
fontweight='bold', alpha=0.8, y=1.)

for i in range(0, len(numerical_features)):
    plt.subplot(4, 2, i+1)
    sns.kdeplot(x=df[numerical_features[i]], shade=True,
hue=df['income'], color='r')
    plt.xlabel(numerical_features[i], fontsize=15)
    plt.tight_layout()
```

## Univariate Analysis of Numerical Features

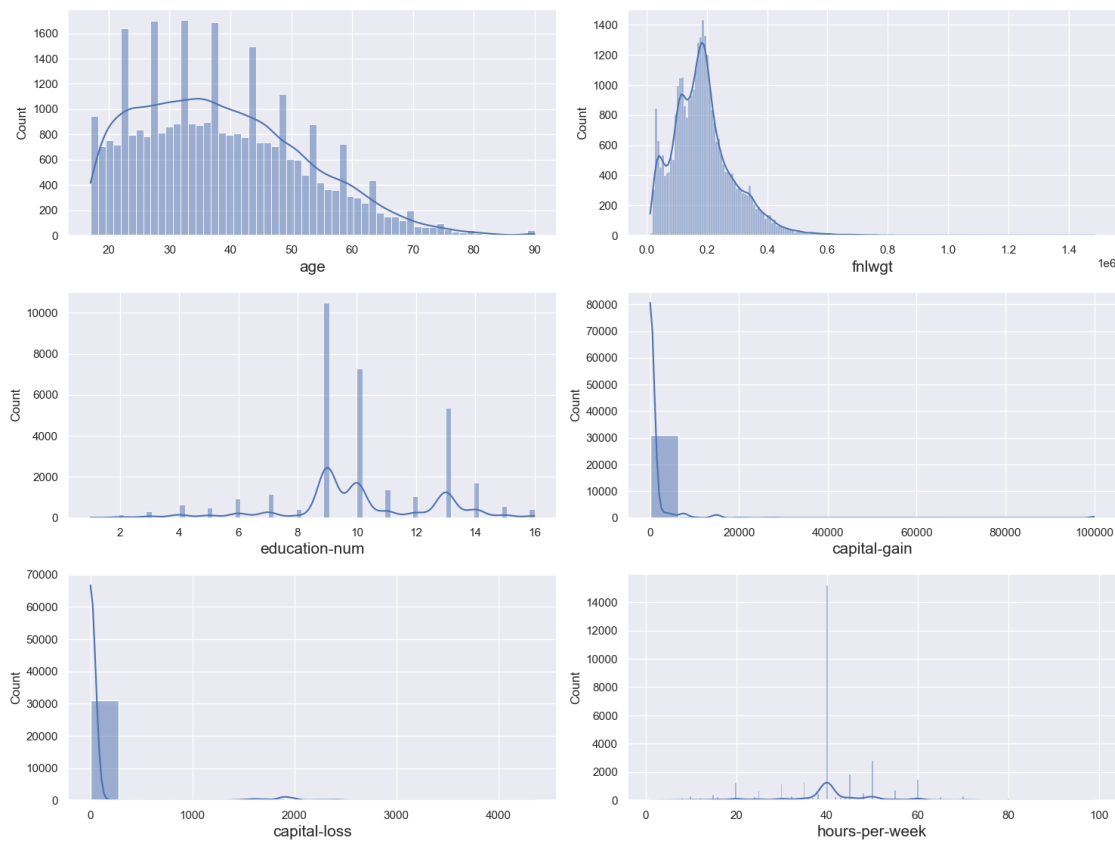


## # Histograms

```
plt.figure(figsize=(15, 15))
plt.suptitle('Univariate Analysis of Numerical Features', fontsize=20,
fontweight='bold', alpha=0.8, y=1.)

for i in range(0, len(numerical_features)):
    plt.subplot(4, 2, i+1)
    sns.histplot(x=df[numerical_features[i]], kde=True, color='b')
    plt.xlabel(numerical_features[i], fontsize=15)
    plt.tight_layout()
```

## Univariate Analysis of Numerical Features



## Observations

- age, fnlwgt is rightly skewed.
- Generally people work for 30 to 40 hours per week.
- There are outliers in all the columns.










## Categorical Features

```
plt.figure(figsize=(20, 30))
plt.suptitle('Univariate Analysis of Categorical Features',
             fontsize=20, fontweight='bold', alpha=0.8, y=1.)
cat = ['workclass', 'education', 'marital-status', 'occupation',
       'relationship', 'race', 'sex', 'native-country', 'income']

for i in range(0, len(cat)):
    plt.subplot(9, 1, i+1)
    ax = sns.countplot(x=df[cat[i]])
    for p in ax.patches:
        ax.annotate('{:.1f}'.format(p.get_height()), (p.get_x()+0.25,
p.get_height()+0.01))
    plt.xlabel(cat[i], fontsize=15)
    plt.xticks(rotation=90)
    plt.tight_layout()
```

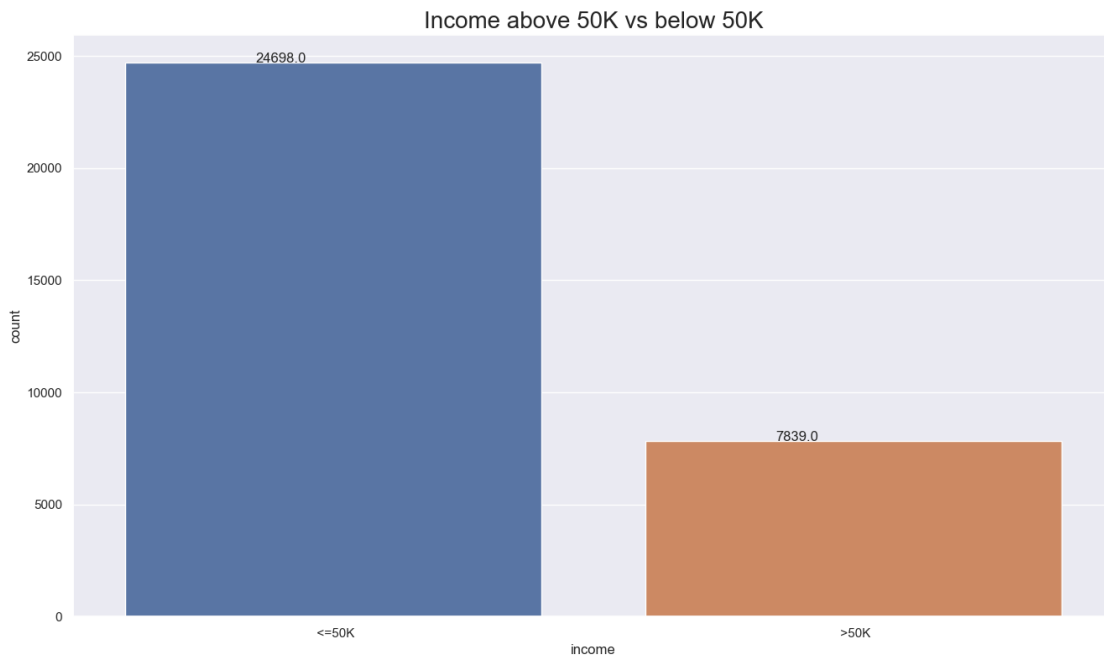


Univariate Analysis of Categorical Features



### Countplot of income column

```
plt.figure(figsize=(16, 9))
plt.title("Income above 50K vs below 50K", fontsize=20)
ax = sns.countplot(x='income', data=df)
for p in ax.patches:
    ax.annotate('{:.1f}'.format(p.get_height()), (p.get_x()+0.25,
p.get_height()+0.01))
plt.show()
```



### Observation:

- Most people earn less than 50K.

### 3.3.2 Biivariate Analysis

#### Categorical features

# Creating a dataframe of categorical columns

```
df_categoric = df[categorical_features]
df_categoric.head()
```

	workclass	education	marital-status	
occupation \				
0	State-gov	Bachelors	Never-married	Adm-
clerical				
1	Self-emp-not-inc	Bachelors	Married-civ-spouse	Exec-
managerial				
2	Private	HS-grad	Divorced	Handlers-
cleaners				
3	Private	11th	Married-civ-spouse	Handlers-

```

cleaners
4          Private  Bachelors  Married-civ-spouse  Prof-
specialty

```

```

      relationship  race    sex  native-country  income
0  Not-in-family  White  Male  United-States  <=50K
1      Husband  White  Male  United-States  <=50K
2  Not-in-family  White  Male  United-States  <=50K
3      Husband  Black  Male  United-States  <=50K
4      Wife  Black  Female  Cuba  <=50K

```

```

plt.figure(figsize=(30, 40))
plt.suptitle('Count of people earning more or less than 50K based on
other categorical features',
            fontsize=20, fontweight='bold', alpha=0.8, y=1.)

```

```

column_names = df_categoric.columns

for i in range(0, len(column_names)):
    plt.subplot(9, 1, i+1)
    ax = sns.countplot(x=df_categoric[column_names[i]],
hue=df_categoric['income'])
    for p in ax.patches:
        ax.annotate('{:.1f}'.format(p.get_height()), (p.get_x()+0.25,
p.get_height()+0.01))
    plt.xlabel(column_names[i], fontsize=15)
    plt.tight_layout()
plt.show()

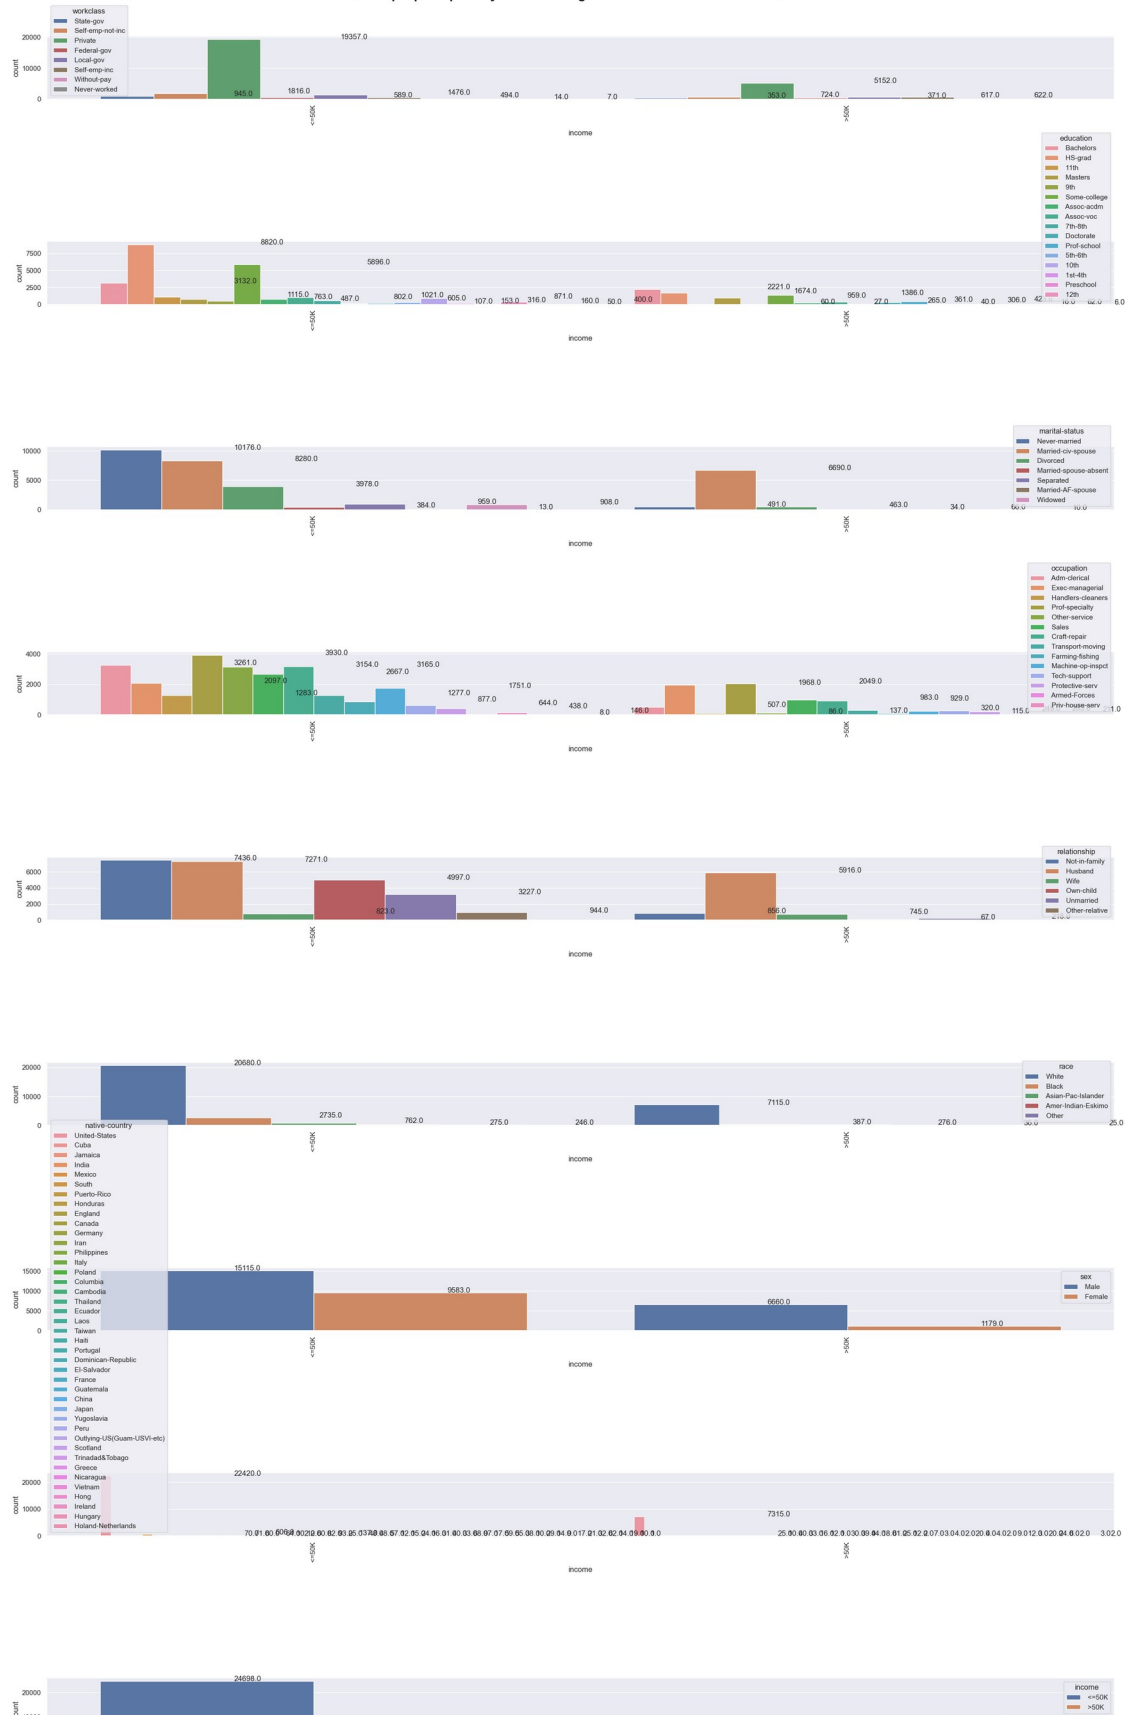
```



```
column_names = df_categoric.columns

for i in range(0, len(column_names)):
    plt.subplot(9, 1, i+1)
    ax = sns.countplot(x='income', hue=column_names[i],
data=df_categoric)
    for p in ax.patches:
        ax.annotate('{:.1f}'.format(p.get_height()), (p.get_x()+0.25,
p.get_height()+0.01))
    plt.xticks(rotation = 90)
    plt.tight_layout()
```

Count of people seperately for each categorical feature based on their income



### Numerical Features

*# Creating a dataframe leaving the two columns 'Classes' and 'Region'*

```
df_numeric = df[numerical_features]
df_numeric.head()
```

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week
0	39	77516	13	2174	0	
1	50	83311	13	0	0	
2	38	215646	9	0	0	
3	53	234721	7	0	0	
4	28	338409	13	0	0	

*# Adding the categorical column 'income' to the numeric dataframe*

```
df_numeric['income'] = df_categorical['income']
df_numeric.head()
```

	age	fnlwgt	education-num	capital-gain	capital-loss	hours-per-week
0	39	77516	13	2174	0	
1	50	83311	13	0	0	
2	38	215646	9	0	0	
3	53	234721	7	0	0	
4	28	338409	13	0	0	

	income
0	<=50K
1	<=50K
2	<=50K
3	<=50K
4	<=50K

### Creating a numeric column on the basis of the categorical column income

*# In this column the value '<=50K' will be represented as '0' and '>50K' as '1'*

```
df_numeric['earning_class'] = df_categorical['income'].map({'<=50K':0,
```

```
' >50K':1})
df_numeric.head()

   age  fnlwgt  education-num  capital-gain  capital-loss  hours-per-
week \
0   39   77516             13           2174             0
40
1   50   83311             13              0             0
13
2   38  215646              9              0             0
40
3   53  234721              7              0             0
40
4   28  338409             13              0             0
40

   income  earning_class
0  <=50K             0
1  <=50K             0
2  <=50K             0
3  <=50K             0
4  <=50K             0
```

*# Now dropping the categorical class 'income'*

```
df_numeric.drop(columns=['income'], axis=1, inplace=True)
df_numeric.head()

   age  fnlwgt  education-num  capital-gain  capital-loss  hours-per-
week \
0   39   77516             13           2174             0
40
1   50   83311             13              0             0
13
2   38  215646              9              0             0
40
3   53  234721              7              0             0
40
4   28  338409             13              0             0
40

   earning_class
0             0
1             0
2             0
3             0
4             0
```

```
plt.figure(figsize=(20, 20))
plt.suptitle('Relation of earning_class feature with other numeric
features', fontsize=20, fontweight='bold', alpha=0.8, y=1.)
```

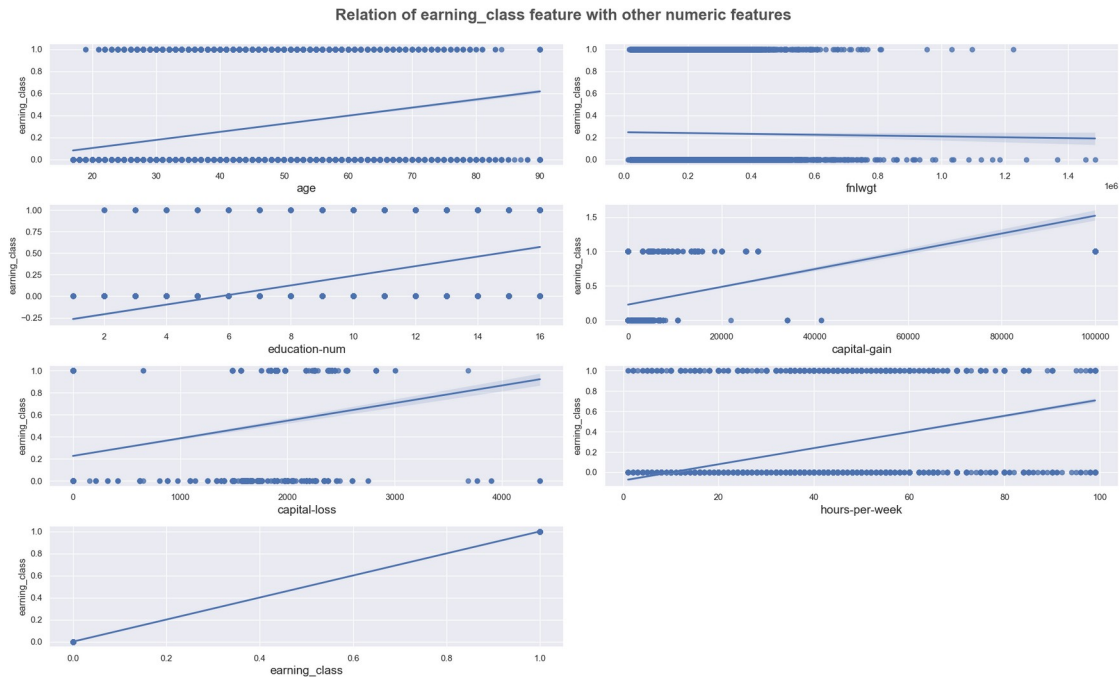


```

column_names = df_numeric.columns

for i in range(0, len(column_names)):
    plt.subplot(7, 2, i+1)
    sns.regplot(x=df_numeric[column_names[i]],
y=df_numeric['earning_class'])
    plt.xlabel(column_names[i], fontsize=15)
    plt.tight_layout()
plt.show()

```



## Transforming income from categorical to numeric in the original dataset

```

df['income'] = df['income'].apply(lambda x:x.replace("<=50K", "0"))
df['income'] = df['income'].apply(lambda x:x.replace(">50K", "1"))
df['income'] = df['income'].astype(int)
df.head()

```

	age	workclass	fnlwgt	education	education-num	\
0	39	State-gov	77516	Bachelors	13	
1	50	Self-emp-not-inc	83311	Bachelors	13	
2	38	Private	215646	HS-grad	9	
3	53	Private	234721	11th	7	
4	28	Private	338409	Bachelors	13	

	sex	marital-status	occupation	relationship	race
0	Male	Never-married	Adm-clerical	Not-in-family	White
1		Married-civ-spouse	Exec-managerial	Husband	White

```

Male
2          Divorced   Handlers-cleaners   Not-in-family   White
Male
3   Married-civ-spouse   Handlers-cleaners           Husband   Black
Male
4   Married-civ-spouse           Prof-specialty           Wife   Black
Female

```

```

      capital-gain  capital-loss  hours-per-week  native-country  income
0           2174           0           40   United-States           0
1              0           0           13   United-States           0
2              0           0           40   United-States           0
3              0           0           40   United-States           0
4              0           0           40           Cuba           0

```

*# Checking numeric features again*

```

numeric_features = [feature for feature in df.columns if
df[feature].dtype != '0']
print(f"The number of Numerical features are: {len(numeric_features)},
and the column names are:\n{numeric_features}")

```

The number of Numerical features are: 7, and the column names are:  
['age', 'fnlwgt', 'education-num', 'capital-gain', 'capital-loss',  
'hours-per-week', 'income']

### 3.3.3 Multivariate Analysis

#### Checking Multicollinearity in the numerical features

```
df[numeric_features].corr()
```

```

      capital-loss \
age          1.000000 -0.076447      0.036224      0.077676
0.057745
fnlwgt       -0.076447  1.000000     -0.043388      0.000429      -
0.010260
education-num  0.036224 -0.043388      1.000000      0.122664
0.079892
capital-gain   0.077676  0.000429      0.122664      1.000000      -
0.031639
capital-loss   0.057745 -0.010260      0.079892     -0.031639
1.000000
hours-per-week  0.068515 -0.018898      0.148422      0.078408
0.054229

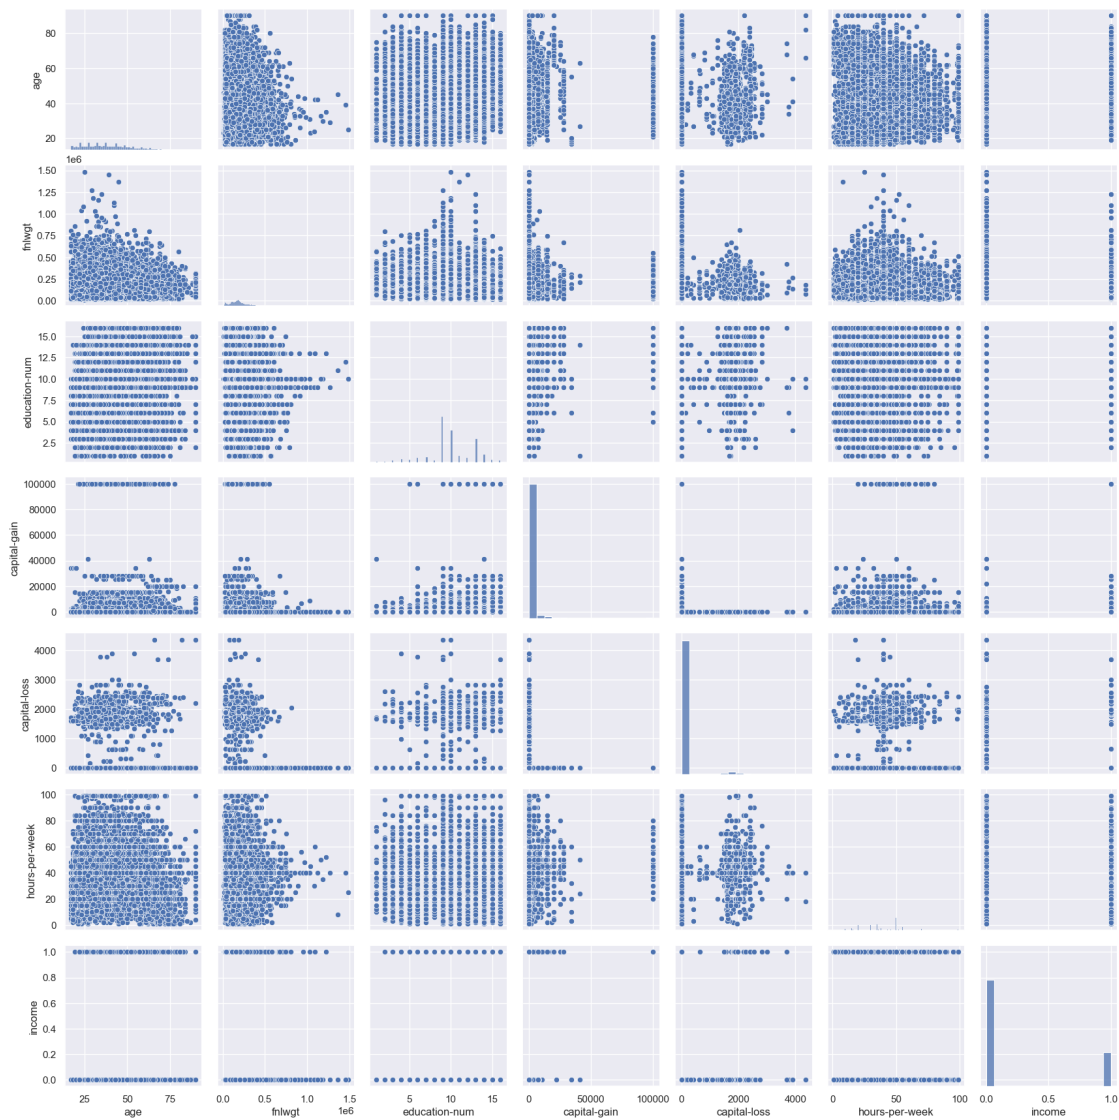
```

income	0.234037	-0.009502	0.335272	0.223336
0.150501				

	hours-per-week	income
age	0.068515	0.234037
fnlwgt	-0.018898	-0.009502
education-num	0.148422	0.335272
capital-gain	0.078408	0.223336
capital-loss	0.054229	0.150501
hours-per-week	1.000000	0.229658
income	0.229658	1.000000

### Graphical Representation

```
sns.pairplot(df[numeric_features])
plt.show()
```



```
plt.figure(figsize = (15,10))
sns.heatmap(df[numeric_features].corr(), cmap="CMRmap", annot=True)
plt.show()
```



## Checking Multicollinearity in the categorical features

- Using Chi-squared Test

categorical\_features

```
['workclass',
 'education',
 'marital-status',
 'occupation',
 'relationship',
 'race',
 'sex',
 'native-country',
 'income']
```

```
from scipy.stats import chi2_contingency
```

```
chi2_test = []
```

```
for feature in categorical_features:
    if chi2_contingency(pd.crosstab(df['income'], df[feature]))[1] <
0.05:
```

```

        chi2_test.append('Reject Null Hypothesis')
    else:
        chi2_test.append('Fail to Reject Null Hypothesis')

result = pd.DataFrame(data=[categorical_features, chi2_test]).T
result.columns = ['Column', 'Hypothesis Result']
result

```

	Column	Hypothesis Result
0	workclass	Reject Null Hypothesis
1	education	Reject Null Hypothesis
2	marital-status	Reject Null Hypothesis
3	occupation	Reject Null Hypothesis
4	relationship	Reject Null Hypothesis
5	race	Reject Null Hypothesis
6	sex	Reject Null Hypothesis
7	native-country	Reject Null Hypothesis
8	income	Reject Null Hypothesis

### Observations:

- Distribution of Numerical columns age, fnlwgt are rightly skewed.
- In the columns capital-gain and capital-loss most of the values are **powerlaw distributed**.
- Most of the employees work in the **Private** sector.
- Most of the people participated in the census are **white, male, Married-civ-spouse** and their native country is **USA**.
- All Numerical features have outliers.
- Also we can see the dataset is imbalanced as the target column income has 75% datapoints in the **<50K** class, so we need to balance it.

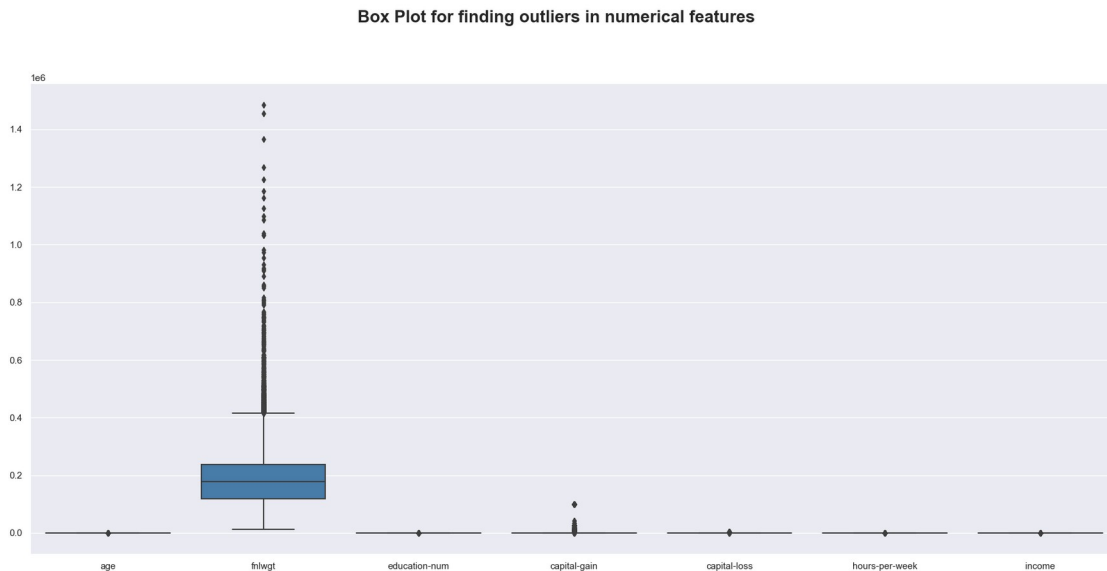
## 4. Data Pre-Processing

### Checking Outliers

```

fig = plt.figure( figsize=(23, 10))
plt.suptitle('Box Plot for finding outliers in numerical features',
fontsize=20, fontweight='bold', alpha=1, y=1)
stud_bplt = sns.boxplot(orient='v', data=df[numeric_features],
palette="Set1")
stud_bplt.plot()
plt.show()

```



**There are outliers**

**Checking for multicollinearity in numeric features**

*# importing library*

```
from statsmodels.stats.outliers_influence import
variance_inflation_factor
```

*# Creating function*

```
def calc_vif(X):
    vif = pd.DataFrame()
    vif["variables"] = X.columns
    vif["VIF"] = [variance_inflation_factor(X.values, i) for i in
range(X.shape[1])]

    return(vif)
```

```
df1 = df[numeric_features]
calc_vif(df1)
```

	variables	VIF
0	age	7.293626
1	fnlwgt	3.716829
2	education-num	11.205498
3	capital-gain	1.081154
4	capital-loss	1.078271
5	hours-per-week	9.776523
6	income	1.549287

**Observation:**

- Only education-num has a value of >10, so it has multicollinearity.

```
# Deleting 'education num'
```

```
df = df.drop(['education-num'], axis = 1)
df.head()
```

	age	workclass	fnlwgt	education	marital-status \
0	39	State-gov	77516	Bachelors	Never-married
1	50	Self-emp-not-inc	83311	Bachelors	Married-civ-spouse
2	38	Private	215646	HS-grad	Divorced
3	53	Private	234721	11th	Married-civ-spouse
4	28	Private	338409	Bachelors	Married-civ-spouse

	gain \	occupation	relationship	race	sex	capital-
0		Adm-clerical	Not-in-family	White	Male	2174
1		Exec-managerial	Husband	White	Male	0
2		Handlers-cleaners	Not-in-family	White	Male	0
3		Handlers-cleaners	Husband	Black	Male	0
4		Prof-specialty	Wife	Black	Female	0

	capital-loss	hours-per-week	native-country	income
0	0	40	United-States	0
1	0	13	United-States	0
2	0	40	United-States	0
3	0	40	United-States	0
4	0	40	Cuba	0

**Storing this data 1st in folder then in MongoDB for later use**

```
try:
    df.to_csv("dataset/adult_processed.csv", index=None)
except Exception as err:
    print("Error is: ", err)
else:
    print("Processed csv file created successfully.")
```

Processed csv file created successfully.

**MongoDB part**

```
income_dict = df.to_dict('records')
```

```
# connecting with the server
```

```
try:
    client =
```

```

pymongo.MongoClient("mongodb+srv://ineuron:Project1@cluster0.rp4qzrr.m
ongodb.net/?retryWrites=true&w=majority")
    db = client
except Exception as e:
    print(e)
else:
    print("Connection to MongoDB server is successful.")
finally:
    print(db)

```

Connection to MongoDB server is successful.

```

MongoClient(host=['ac-rnik5cy-shard-00-00.rp4qzrr.mongodb.net:27017',
'ac-rnik5cy-shard-00-02.rp4qzrr.mongodb.net:27017', 'ac-rnik5cy-shard-
00-01.rp4qzrr.mongodb.net:27017'], document_class=dict,
tz_aware=False, connect=True, retrywrites=True, w='majority',
authsource='admin', replicaset='atlas-jltif8-shard-0', tls=True)

```

*# Creating database and collection*

```

database = client["income_census_data"]
collection = database['income_census']

try:
    collection.insert_many(income_dict)
except Exception as e:
    print(e)
else:
    print("Records inserted successfully.")

```

Records inserted successfully.

### Loading the data from MongoDB

```

db = client.income_census_data
collect_names = db.list_collection_names()
collect_names

['income_census']

final_df = pd.DataFrame(list(db.income_census.find()))
final_df

```

	_id	age	workclass	fnlwgt
education \				
0	63651ee3a1b4514aba84746d	39	State-gov	77516
Bachelors				
1	63651ee3a1b4514aba84746e	50	Self-emp-not-inc	83311
Bachelors				
2	63651ee3a1b4514aba84746f	38	Private	215646
HS-grad				
3	63651ee3a1b4514aba847470	53	Private	234721
11th				



4	63651ee3a1b4514aba847471	28	Private	338409
Bachelors				
...	...	...	...	...
...				
32532	63651ee4a1b4514aba84f381	27	Private	257302
Assoc-acdm				
32533	63651ee4a1b4514aba84f382	40	Private	154374
HS-grad				
32534	63651ee4a1b4514aba84f383	58	Private	151910
HS-grad				
32535	63651ee4a1b4514aba84f384	22	Private	201490
HS-grad				
32536	63651ee4a1b4514aba84f385	52	Self-emp-inc	287927
HS-grad				

	marital-status	occupation	relationship	race
\				
0	Never-married	Adm-clerical	Not-in-family	White
1	Married-civ-spouse	Exec-managerial	Husband	White
2	Divorced	Handlers-cleaners	Not-in-family	White
3	Married-civ-spouse	Handlers-cleaners	Husband	Black
4	Married-civ-spouse	Prof-specialty	Wife	Black
...	...	...	...	...
32532	Married-civ-spouse	Tech-support	Wife	White
32533	Married-civ-spouse	Machine-op-inspct	Husband	White
32534	Widowed	Adm-clerical	Unmarried	White
32535	Never-married	Adm-clerical	Own-child	White
32536	Married-civ-spouse	Exec-managerial	Wife	White

country	sex	capital-gain	capital-loss	hours-per-week	native-
\					
0	Male	2174	0	40	United-
States					
1	Male	0	0	13	United-
States					
2	Male	0	0	40	United-
States					
3	Male	0	0	40	United-

States					
4	Female	0	0	40	
Cuba					
...	...	...	...	...	
...					
32532	Female	0	0	38	United-
States					
32533	Male	0	0	40	United-
States					
32534	Female	0	0	40	United-
States					
32535	Male	0	0	20	United-
States					
32536	Female	15024	0	40	United-
States					

	income
0	0
1	0
2	0
3	0
4	0
...	...
32532	0
32533	1
32534	0
32535	0
32536	1

[32537 rows x 15 columns]

*# Dropping the '\_id' column*

```
final_df.drop(['_id'], axis = 1, inplace=True)
final_df
```

	age	workclass	fnlwgt	education	marital-
status \					
0	39	State-gov	77516	Bachelors	Never-
married					
1	50	Self-emp-not-inc	83311	Bachelors	Married-civ-
spouse					
2	38	Private	215646	HS-grad	
Divorced					
3	53	Private	234721	11th	Married-civ-
spouse					
4	28	Private	338409	Bachelors	Married-civ-
spouse					
...	...	...	...	...	..
.					

32532	27	Private	257302	Assoc-acdm	Married-civ-spouse
32533	40	Private	154374	HS-grad	Married-civ-spouse
32534	58	Private	151910	HS-grad	Widowed
32535	22	Private	201490	HS-grad	Never-married
32536	52	Self-emp-inc	287927	HS-grad	Married-civ-spouse

gain \	occupation	relationship	race	sex	capital-
0	Adm-clerical	Not-in-family	White	Male	
2174					
1	Exec-managerial	Husband	White	Male	
0					
2	Handlers-cleaners	Not-in-family	White	Male	
0					
3	Handlers-cleaners	Husband	Black	Male	
0					
4	Prof-specialty	Wife	Black	Female	
0					
...	...	...	...	...	.
..					
32532	Tech-support	Wife	White	Female	
0					
32533	Machine-op-inspct	Husband	White	Male	
0					
32534	Adm-clerical	Unmarried	White	Female	
0					
32535	Adm-clerical	Own-child	White	Male	
0					
32536	Exec-managerial	Wife	White	Female	
15024					

	capital-loss	hours-per-week	native-country	income
0	0	40	United-States	0
1	0	13	United-States	0
2	0	40	United-States	0
3	0	40	United-States	0
4	0	40	Cuba	0
...	...	...	...	...
32532	0	38	United-States	0
32533	0	40	United-States	1
32534	0	40	United-States	0
32535	0	20	United-States	0
32536	0	40	United-States	1

[32537 rows x 14 columns]

## Creating independent and dependent variables

```
X = df.drop('income', axis = 1)
y = df['income']
```

```
# Doing Test Train split
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.33, random_state=42)
```

```
X_train.shape
```

```
(21799, 13)
```

```
X_test.shape
```

```
(10738, 13)
```

```
category_var = [col for col in X.columns if X[col].dtypes == object]
category_var
```

```
['workclass',
 'education',
 'marital-status',
 'occupation',
 'relationship',
 'race',
 'sex',
 'native-country']
```

```
numeric_var = [col for col in X.columns if X[col].dtypes != object]
numeric_var
```

```
['age', 'fnlwgt', 'capital-gain', 'capital-loss', 'hours-per-week']
```

## Encoding

### To do encoding

- <https://www.kaggle.com/code/subinium/11-categorical-encoders-and-benchmark>

```
# importing library
```

```
import category_encoders as ce
```

```
one_hot = ce.OneHotEncoder(cols = category_var, handle_unknown =
'ignore')
```

```
# Creating dataframe for categorical variables which converted to one
hot encoded variables.
```

```
X_train_one_hot = pd.DataFrame(one_hot.fit_transform(X_train))
```

```
X_test_one_hot = pd.DataFrame(one_hot.transform(X_test))
```

```

X_train_one_hot.index = X_train.index
X_test_one_hot.index = X_test.index

num_X_train = X_train[numeric_var]
num_X_test = X_test[numeric_var]

# Joining numerical and one hot encoded variables to create our final
X_train and X_test.
X_train_new = pd.concat([num_X_train, X_train_one_hot], axis = 1)
X_test_new = pd.concat([num_X_test, X_test_one_hot], axis = 1)

```

### Scaling the data

```
from sklearn.preprocessing import StandardScaler
```

```

scaler = StandardScaler()
scaler

StandardScaler()

scale = scaler.fit(X_train_new)
scale

```

```
StandardScaler()
```

### *# Printing the mean*

```

print(scale.mean_)

[3.85290151e+01 1.90201068e+05 1.02077577e+03 9.01778063e+01
 4.04248360e+01 3.85290151e+01 7.54392403e-01 6.33515299e-02
 7.74347447e-02 3.46346163e-02 2.96802606e-02 3.99559613e-02
 3.66989311e-04 1.83494656e-04 1.90201068e+05 2.25102069e-01
 3.22858847e-01 3.26620487e-02 2.78911877e-02 1.01380797e-02
 1.58722877e-02 1.64778201e-01 1.36244782e-02 3.71576678e-02
 5.24794715e-02 5.00022937e-03 1.27528786e-02 1.70191293e-02
 4.20202762e-02 1.90375705e-02 1.60557824e-03 3.29969265e-01
 4.57681545e-01 1.35464930e-01 3.10105968e-02 3.22033121e-02
 1.28446259e-02 8.25725951e-04 4.21120235e-02 1.16335612e-01
 1.96798018e-02 1.00371577e-01 4.77544842e-02 1.12849213e-01
 1.22895546e-01 1.84595624e-01 1.26978302e-01 2.81664297e-02
 6.20211936e-02 3.10105968e-02 4.95435570e-03 2.75241984e-04
 1.58906372e-01 4.92683151e-02 4.01899170e-01 2.53130878e-01
 1.05968164e-01 3.08271022e-02 8.52974907e-01 9.59677049e-02
 3.27996697e-02 9.63346943e-03 8.62424882e-03 6.68195789e-01
 3.31804211e-01 1.02077577e+03 9.01778063e+01 4.04248360e+01
 9.11417955e-01 2.00009175e-02 4.49561907e-03 3.44052479e-03
 1.65145190e-03 2.24780953e-03 2.47717785e-03 1.78907289e-03
 6.65168127e-03 3.66989311e-03 3.48639846e-03 2.01844121e-03
 1.37620992e-03 3.80751411e-03 4.12862975e-04 2.38543052e-03

```

```

1.28446259e-03 2.24780953e-03 2.01844121e-03 7.33978623e-04
5.96357631e-04 1.05509427e-03 1.51383091e-03 1.10096793e-03
2.93591449e-03 2.15606220e-03 4.58736639e-04 2.70654617e-03
1.97256755e-03 4.58736639e-04 2.11018854e-03 7.33978623e-04
5.04610303e-04 6.88104959e-04 7.79852287e-04 3.66989311e-04
3.66989311e-04 5.96357631e-04 9.17473279e-04 3.21115648e-04
4.58736639e-05]

```

### **Saving the scale to use it later to transform the data and predict the values**

*# To save a Standard scaler object*

```
import pickle
```

```
with open('scaled.pkl', 'wb') as f:
    pickle.dump(scale, f)
```

*# Loading the scaled object to transform the data*

```
with open('scaled.pkl', 'rb') as f:
    scaled = pickle.load(f)
```

*# Now transforming the train and test dataset*

```
X_train_tf = scaled.transform(X_train_new)
X_test_tf = scaled.transform(X_test_new)
```

*# checking the transformed data*

X\_train\_tf

```

array([[ -1.36146381,  -0.76094509,  -0.14229344, ...,  -0.03030373,
        -0.01792258,  -0.00677317],
       [  0.62242593,  -0.2277653 ,  -0.14229344, ...,  -0.03030373,
        -0.01792258,  -0.00677317],
       [ -0.25930285,   1.04760296,   0.87502844, ...,  -0.03030373,
        -0.01792258,  -0.00677317],
       ...,
       [ -1.50841861,   0.24732597,  -0.14229344, ...,  -0.03030373,
        -0.01792258,  -0.00677317],
       [ -0.33278024,   0.50027559,  -0.14229344, ...,  -0.03030373,
        -0.01792258,  -0.00677317],
       [ -1.14103162,   1.3249031 ,  -0.14229344, ...,  -0.03030373,
        -0.01792258,  -0.00677317]])

```

X\_test\_tf

```

array([[ -0.55321244,  -1.48770587,  -0.14229344, ...,  -0.03030373,
        -0.01792258,  -0.00677317],
       [ -1.06755422,  -0.56578056,  -0.14229344, ...,  -0.03030373,
        -0.01792258,  -0.00677317],
       [  1.5776321 ,  -0.87095864,  -0.14229344, ...,  -0.03030373,

```

```

-0.01792258, -0.00677317],
...,
[-0.84712203, -0.0192497, -0.14229344, ..., -0.03030373,
-0.01792258, -0.00677317],
[-0.03887065, 0.14500256, -0.14229344, ..., -0.03030373,
-0.01792258, -0.00677317],
[-0.62668984, 1.42913245, -0.14229344, ..., -0.03030373,
-0.01792258, -0.00677317]])

```

## 5. Model Building

```

from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.metrics import accuracy_score, confusion_matrix,
classification_report, roc_curve, roc_auc_score

```

### Logistic Regression

```

lgr = LogisticRegression()
lgr.fit(X_train_new, y_train)
lgr_pred = lgr.predict(X_test_new)
ac_lr = accuracy_score(y_test, lgr_pred)
roc_logr = roc_auc_score(y_test, lgr_pred)
print('Logistic Regression accuracy score:{0:0.2f}%'.
format(ac_lr*100))
print('Logistic Regression ROC score:{0:0.2f}%'. format(roc_logr*100))

```

```

Logistic Regression accuracy score:79.74%
Logistic Regression ROC score:62.02%

```

### SVC

```

classifier = SVC(random_state = 0, kernel = 'rbf')
classifier.fit(X_train_new, y_train)
y_pred = classifier.predict(X_test_new)
ac_svc = accuracy_score(y_test,y_pred)
roc_svc = roc_auc_score(y_test,y_pred)
print('SVC accuracy score:{0:0.2f}%'. format(ac_svc*100))
print('SVC ROC score:{0:0.2f}%'. format(roc_svc*100))

```

```

SVC accuracy score:79.55%
SVC ROC score:57.83%

```

### # Saving the models

```

with open('lgr.pkl', 'wb') as f:
    pickle.dump(lgr, f)

with open('svc.pkl', 'wb') as f:
    pickle.dump(classifier, f)

```

### # Loading the models

```
with open('lgr.pkl', 'rb') as f:
    clf_logreg = pickle.load(f)
```

```
with open('svc.pkl', 'rb') as f:
    clf_SVC = pickle.load(f)
```

### Evaluation of the models

```
models = {clf_logreg: 'LogisticRegression',
          clf_SVC: 'SVC',
          }
```

```
def train(algo, name, X_train, y_train, X_test, y_test):
    algo.fit(X_train, y_train)
    y_pred = algo.predict(X_test)
    score = accuracy_score(y_test, y_pred)
    print(f"-----")
    {name}-----")
    print(f"Accuracy Score for {name}: {score*100:.4f}%")
    return y_test, y_pred, score
```

*# acc\_res function calculates confusion matrix*

```
def acc_res(y_test, y_pred):
    null_accuracy = y_test.value_counts()[0]/len(y_test)
    print(f"Null Accuracy: {null_accuracy*100:.4f}%")
    print("Confusion Matrix")
    matrix = confusion_matrix(y_test, y_pred)
    print(matrix)
    print("+++++")
    TN = matrix[0,0]
    FP = matrix[0,1]
    FN = matrix[1,0]
    TP = matrix[1,1]
    accuracy_score = (TN+TP) / float(TP+TN+FP+FN)
    recall_score = (TP) / float(TP+FN)
    specificity = TN / float(TN+FP)
    FPR = FP / float(FP+TN)
    precision_score = TP / float(TP+FP)
    print(f"Accuracy Score: {accuracy_score*100:.4f}%")
    print(f"Recall Score: {recall_score*100:.4f}%")
    print(f"Specificity Score: {specificity*100:.4f}%")
    print(f"False Positive Rate: {FPR*100:.4f}%")
    print(f"Precision Score: {precision_score*100:.4f}%")
    print("+++++")
    print("Classification Report")
    print(classification_report(y_test, y_pred))
```

```
def main(models):
    accuracy_scores = []
    for algo, name in models.items():
```



```

        y_test_train,y_pred,acc_score =
train(algo,name,X_train_new,y_train,X_test_new,y_test)
        acc_res(y_test_train,y_pred)
        accuracy_scores.append(acc_score)
    return accuracy_scores

```

```
accuracy_scores = main(models)
```

```
-----
LogisticRegression-----
```

Accuracy Score for LogisticRegression: 79.7448%

Null Accuracy: 75.7962%

Confusion Matrix

```
[[7844  295]
 [1880  719]]
```

```
+++++
```

Accuracy Score: 79.7448%

Recall Score: 27.6645%

Specificity Score: 96.3755%

False Positive Rate: 3.6245%

Precision Score: 70.9073%

```
+++++
```

Classification Report

	precision	recall	f1-score	support
0	0.81	0.96	0.88	8139
1	0.71	0.28	0.40	2599
accuracy			0.80	10738
macro avg	0.76	0.62	0.64	10738
weighted avg	0.78	0.80	0.76	10738

```
-----
SVC-----
```

Accuracy Score for SVC: 79.5493%

Null Accuracy: 75.7962%

Confusion Matrix

```
[[8133    6]
 [2190   409]]
```

```
+++++
```

Accuracy Score: 79.5493%

Recall Score: 15.7368%

Specificity Score: 99.9263%

False Positive Rate: 0.0737%

Precision Score: 98.5542%

```
+++++
```

Classification Report				
	precision	recall	f1-score	support
0	0.79	1.00	0.88	8139
1	0.99	0.16	0.27	2599
accuracy			0.80	10738
macro avg	0.89	0.58	0.58	10738
weighted avg	0.84	0.80	0.73	10738

**As both the models are giving almost same accuracy so we go with logistic regression**

## Hyperparameter Tuning using GridSearchCV

```
from sklearn.model_selection import GridSearchCV
```

## Logistic Regression

```
parameters = {
    'penalty' : ['l1', 'l2'],
    'C'       : np.logspace(-3, 3, 7),
    'solver'  : ['newton-cg', 'lbfgs', 'liblinear'],
}

clf = GridSearchCV(clf_logreg,                    # model
                   param_grid = parameters,       # hyperparameters
                   scoring='accuracy',            # metric for scoring
                   verbose=2,
                   cv=10)

clf.fit(X_train_new, y_train)
print("Tuned Hyperparameters :", clf.best_params_)
print("Accuracy :", clf.best_score_)
```

[illegible]

```
[CV] END .....C=0.001, penalty=l1, solver=newton-cg; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l2, solver=newton-cg; total
time= 6.3s
[CV] END .....C=0.001, penalty=l2, solver=newton-cg; total
time= 6.7s
[CV] END .....C=0.001, penalty=l2, solver=newton-cg; total
time= 8.2s
[CV] END .....C=0.001, penalty=l2, solver=newton-cg; total
time= 5.3s
```

```
[CV] END .....C=0.001, penalty=l2, solver=newton-cg; total
time= 5.6s
[CV] END .....C=0.001, penalty=l2, solver=newton-cg; total
time= 5.5s
[CV] END .....C=0.001, penalty=l2, solver=newton-cg; total
time= 5.4s
[CV] END .....C=0.001, penalty=l2, solver=newton-cg; total
time= 5.5s
[CV] END .....C=0.001, penalty=l2, solver=newton-cg; total
time= 6.1s
[CV] END .....C=0.001, penalty=l2, solver=newton-cg; total
time= 5.0s
[CV] END .....C=0.001, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=0.001, penalty=l2, solver=lbfgs; total
time= 0.4s
[CV] END .....C=0.001, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.001, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.001, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END .....C=0.001, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.001, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.001, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.001, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=0.001, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.001, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.001, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.001, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.001, penalty=l2, solver=liblinear; total
time= 0.1s
```

[illegible]

```

[CV] END .....C=0.01, penalty=l1, solver=liblinear; total
time= 0.2s
[CV] END .....C=0.01, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.01, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.01, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.01, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.01, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.01, penalty=l2, solver=newton-cg; total
time= 5.1s
[CV] END .....C=0.01, penalty=l2, solver=newton-cg; total
time= 5.0s
[CV] END .....C=0.01, penalty=l2, solver=newton-cg; total
time= 5.5s
[CV] END .....C=0.01, penalty=l2, solver=newton-cg; total
time= 4.6s
[CV] END .....C=0.01, penalty=l2, solver=newton-cg; total
time= 5.5s
[CV] END .....C=0.01, penalty=l2, solver=newton-cg; total
time= 5.0s
[CV] END .....C=0.01, penalty=l2, solver=newton-cg; total
time= 5.3s
[CV] END .....C=0.01, penalty=l2, solver=newton-cg; total
time= 4.8s
[CV] END .....C=0.01, penalty=l2, solver=newton-cg; total
time= 4.7s
[CV] END .....C=0.01, penalty=l2, solver=newton-cg; total
time= 4.8s
[CV] END .....C=0.01, penalty=l2, solver=lbfgs; total
time= 0.2s

[CV] END .....C=0.01, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=0.01, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.01, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.01, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END .....C=0.01, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.01, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.01, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.01, penalty=l2, solver=lbfgs; total

```

[illegible]

```
time= 0.0s
[CV] END .....C=0.1, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.1, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.1, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.1, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.1, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.1, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=0.1, penalty=l1, solver=liblinear; total
time= 0.8s
[CV] END .....C=0.1, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.1, penalty=l1, solver=liblinear; total
time= 0.8s
[CV] END .....C=0.1, penalty=l1, solver=liblinear; total
time= 1.1s
[CV] END .....C=0.1, penalty=l1, solver=liblinear; total
time= 0.2s
[CV] END .....C=0.1, penalty=l1, solver=liblinear; total
time= 0.9s
[CV] END .....C=0.1, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.1, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.1, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.1, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.1, penalty=l2, solver=newton-cg; total
time= 4.6s
[CV] END .....C=0.1, penalty=l2, solver=newton-cg; total
time= 5.0s
[CV] END .....C=0.1, penalty=l2, solver=newton-cg; total
time= 4.7s
[CV] END .....C=0.1, penalty=l2, solver=newton-cg; total
time= 4.5s
[CV] END .....C=0.1, penalty=l2, solver=newton-cg; total
time= 5.4s
[CV] END .....C=0.1, penalty=l2, solver=newton-cg; total
time= 4.6s
[CV] END .....C=0.1, penalty=l2, solver=newton-cg; total
time= 4.4s
[CV] END .....C=0.1, penalty=l2, solver=newton-cg; total
time= 4.3s
[CV] END .....C=0.1, penalty=l2, solver=newton-cg; total
```



```
time= 4.8s
[CV] END .....C=0.1, penalty=l2, solver=newton-cg; total
time= 4.4s
[CV] END .....C=0.1, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.1, penalty=l2, solver=lbfgs; total
time= 0.4s
[CV] END .....C=0.1, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=0.1, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.1, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END .....C=0.1, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.1, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.1, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=0.1, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=0.1, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=0.1, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.1, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.1, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.1, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.1, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.1, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.1, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=0.1, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=0.1, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=1.0, penalty=l1, solver=newton-cg; total
time= 0.0s
[CV] END .....C=1.0, penalty=l1, solver=newton-cg; total
time= 0.0s
[CV] END .....C=1.0, penalty=l1, solver=newton-cg; total
time= 0.0s
[CV] END .....C=1.0, penalty=l1, solver=newton-cg; total
```

```
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=newton-cg; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=newton-cg; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=newton-cg; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=newton-cg; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=newton-cg; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=lbfgs; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=lbfgs; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=lbfgs; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=lbfgs; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=lbfgs; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=lbfgs; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=lbfgs; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=lbfgs; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=liblinear; total
time=    0.1s
[CV] END .....C=1.0, penalty=l1, solver=liblinear; total
time=    0.1s
[CV] END .....C=1.0, penalty=l1, solver=liblinear; total
time=    0.1s
[CV] END .....C=1.0, penalty=l1, solver=liblinear; total
time=    0.2s
[CV] END .....C=1.0, penalty=l1, solver=liblinear; total
time=    0.0s
[CV] END .....C=1.0, penalty=l1, solver=liblinear; total
time=    0.1s
[CV] END .....C=1.0, penalty=l1, solver=liblinear; total
time=    0.3s
[CV] END .....C=1.0, penalty=l1, solver=liblinear; total
time=    0.0s
```

```
[CV] END .....C=1.0, penalty=l1, solver=liblinear; total
time= 0.3s
[CV] END .....C=1.0, penalty=l1, solver=liblinear; total
time= 0.5s
[CV] END .....C=1.0, penalty=l2, solver=newton-cg; total
time= 4.8s
[CV] END .....C=1.0, penalty=l2, solver=newton-cg; total
time= 4.8s
[CV] END .....C=1.0, penalty=l2, solver=newton-cg; total
time= 6.1s
[CV] END .....C=1.0, penalty=l2, solver=newton-cg; total
time= 5.3s
[CV] END .....C=1.0, penalty=l2, solver=newton-cg; total
time= 4.7s
[CV] END .....C=1.0, penalty=l2, solver=newton-cg; total
time= 4.4s
[CV] END .....C=1.0, penalty=l2, solver=newton-cg; total
time= 5.0s
[CV] END .....C=1.0, penalty=l2, solver=newton-cg; total
time= 4.9s
[CV] END .....C=1.0, penalty=l2, solver=newton-cg; total
time= 4.5s
[CV] END .....C=1.0, penalty=l2, solver=newton-cg; total
time= 4.7s
[CV] END .....C=1.0, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=1.0, penalty=l2, solver=lbfgs; total
time= 0.5s
[CV] END .....C=1.0, penalty=l2, solver=lbfgs; total
time= 0.4s
[CV] END .....C=1.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=1.0, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END .....C=1.0, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=1.0, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=1.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=1.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=1.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=1.0, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=1.0, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=1.0, penalty=l2, solver=liblinear; total
time= 0.0s
```

[illegible]

```
[CV] END .....C=10.0, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=10.0, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=10.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=10.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=10.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=10.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=10.0, penalty=l1, solver=liblinear; total
time= 0.2s
[CV] END .....C=10.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=10.0, penalty=l1, solver=liblinear; total
time= 0.2s
[CV] END .....C=10.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=10.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=10.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=10.0, penalty=l2, solver=newton-cg; total
time= 4.9s
[CV] END .....C=10.0, penalty=l2, solver=newton-cg; total
time= 4.4s
[CV] END .....C=10.0, penalty=l2, solver=newton-cg; total
time= 4.7s
[CV] END .....C=10.0, penalty=l2, solver=newton-cg; total
time= 4.5s
[CV] END .....C=10.0, penalty=l2, solver=newton-cg; total
time= 4.6s
[CV] END .....C=10.0, penalty=l2, solver=newton-cg; total
time= 5.2s
[CV] END .....C=10.0, penalty=l2, solver=newton-cg; total
time= 5.2s
[CV] END .....C=10.0, penalty=l2, solver=newton-cg; total
time= 5.1s
[CV] END .....C=10.0, penalty=l2, solver=newton-cg; total
time= 5.0s
[CV] END .....C=10.0, penalty=l2, solver=newton-cg; total
time= 5.5s
[CV] END .....C=10.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=10.0, penalty=l2, solver=lbfgs; total
time= 0.6s
[CV] END .....C=10.0, penalty=l2, solver=lbfgs; total
time= 0.3s
```

[illegible]

```

[CV] END .....C=100.0, penalty=l1, solver=newton-cg; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=newton-cg; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=lbfgs; total
time= 0.0s

[CV] END .....C=100.0, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=lbfgs; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=100.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=100.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=100.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=100.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=100.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=100.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=100.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=100.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=100.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=100.0, penalty=l2, solver=newton-cg; total
time= 5.8s
[CV] END .....C=100.0, penalty=l2, solver=newton-cg; total
time= 6.1s
[CV] END .....C=100.0, penalty=l2, solver=newton-cg; total

```

```
time= 4.5s
[CV] END .....C=100.0, penalty=l2, solver=newton-cg; total
time= 5.0s
[CV] END .....C=100.0, penalty=l2, solver=newton-cg; total
time= 4.4s
[CV] END .....C=100.0, penalty=l2, solver=newton-cg; total
time= 4.5s
[CV] END .....C=100.0, penalty=l2, solver=newton-cg; total
time= 5.5s
[CV] END .....C=100.0, penalty=l2, solver=newton-cg; total
time= 5.2s
[CV] END .....C=100.0, penalty=l2, solver=newton-cg; total
time= 4.6s
[CV] END .....C=100.0, penalty=l2, solver=newton-cg; total
time= 4.6s
[CV] END .....C=100.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=100.0, penalty=l2, solver=lbfgs; total
time= 0.5s
[CV] END .....C=100.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=100.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=100.0, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END .....C=100.0, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=100.0, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=100.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=100.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=100.0, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=100.0, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=100.0, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=100.0, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=100.0, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=100.0, penalty=l2, solver=liblinear; total
time= 0.0s
[CV] END .....C=100.0, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=100.0, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=100.0, penalty=l2, solver=liblinear; total
```



[illegible]

```
time= 0.1s
[CV] END .....C=1000.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=1000.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=1000.0, penalty=l1, solver=liblinear; total
time= 0.1s
[CV] END .....C=1000.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=1000.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=1000.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=1000.0, penalty=l1, solver=liblinear; total
time= 0.0s
[CV] END .....C=1000.0, penalty=l2, solver=newton-cg; total
time= 4.9s
[CV] END .....C=1000.0, penalty=l2, solver=newton-cg; total
time= 4.8s
[CV] END .....C=1000.0, penalty=l2, solver=newton-cg; total
time= 4.9s
[CV] END .....C=1000.0, penalty=l2, solver=newton-cg; total
time= 4.5s
[CV] END .....C=1000.0, penalty=l2, solver=newton-cg; total
time= 4.6s
[CV] END .....C=1000.0, penalty=l2, solver=newton-cg; total
time= 5.2s
[CV] END .....C=1000.0, penalty=l2, solver=newton-cg; total
time= 4.5s
[CV] END .....C=1000.0, penalty=l2, solver=newton-cg; total
time= 4.9s
[CV] END .....C=1000.0, penalty=l2, solver=newton-cg; total
time= 5.7s
[CV] END .....C=1000.0, penalty=l2, solver=newton-cg; total
time= 6.0s
[CV] END .....C=1000.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=1000.0, penalty=l2, solver=lbfgs; total
time= 0.6s
[CV] END .....C=1000.0, penalty=l2, solver=lbfgs; total
time= 0.4s
[CV] END .....C=1000.0, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=1000.0, penalty=l2, solver=lbfgs; total
time= 0.1s
[CV] END .....C=1000.0, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=1000.0, penalty=l2, solver=lbfgs; total
time= 0.2s
[CV] END .....C=1000.0, penalty=l2, solver=lbfgs; total
```

```

time= 0.2s
[CV] END .....C=1000.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=1000.0, penalty=l2, solver=lbfgs; total
time= 0.3s
[CV] END .....C=1000.0, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=1000.0, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=1000.0, penalty=l2, solver=liblinear; total
time= 0.2s

[CV] END .....C=1000.0, penalty=l2, solver=liblinear; total
time= 0.2s
[CV] END .....C=1000.0, penalty=l2, solver=liblinear; total
time= 0.3s
[CV] END .....C=1000.0, penalty=l2, solver=liblinear; total
time= 0.2s
[CV] END .....C=1000.0, penalty=l2, solver=liblinear; total
time= 0.1s
[CV] END .....C=1000.0, penalty=l2, solver=liblinear; total
time= 0.2s
[CV] END .....C=1000.0, penalty=l2, solver=liblinear; total
time= 0.4s
[CV] END .....C=1000.0, penalty=l2, solver=liblinear; total
time= 0.1s
Tuned Hyperparameters : {'C': 100.0, 'penalty': 'l2', 'solver':
'newton-cg'}
Accuracy : 0.8506351074266034

```

*# With tuned parameters*

```

model_logr = LogisticRegression(solver="newton-cg", C=100.0,
penalty='l2')
model_logr.fit(X_train_new, y_train)
pred_logr = model_logr.predict(X_test_new)
ac_lr_tuned=accuracy_score(y_test, pred_logr)
roc_logr_tuned =roc_auc_score(y_test,pred_logr)
print('Logistic Regression accuracy score:{0:0.2f}%'.
format(ac_lr_tuned*100))
print('Logistic Regression ROC score:{0:0.2f}%'.
format(roc_logr_tuned*100))

```

```

Logistic Regression accuracy score:85.17%
Logistic Regression ROC score:76.69%

```

#### Roc\_curve for Logistic Regression Model

```

clf_logreg = LogisticRegression(solver="newton-cg", C=100.0,
penalty='l2')

clf_logreg.fit(X_train_new, y_train)

```

```

LogisticRegression(C=100.0, solver='newton-cg')

ytrain_pred = clf_logreg.predict_proba(X_train_new)
print('Logistic train roc-auc: {}'.format(roc_auc_score(y_train,
ytrain_pred[:,1])))
ytest_pred = clf_logreg.predict_proba(X_test_new)
print('Logistic test roc-auc: {}'.format(roc_auc_score(y_test,
ytest_pred[:,1])))

Logistic train roc-auc: 0.9068401203837859
Logistic test roc-auc: 0.9031225965585163

pred=[]
for model in [clf_logreg]:
    pred.append(pd.Series(model.predict_proba(X_test_new)[:,1]))
final_prediction=pd.concat(pred,axis=1).mean(axis=1)
print('Ensemble test roc-auc:
{}'.format(roc_auc_score(y_test,final_prediction)))

Ensemble test roc-auc: 0.9031225965585163

fpr, tpr, thresholds = roc_curve(y_test, final_prediction)
thresholds

array([2.00000000e+00, 1.00000000e+00, 1.00000000e+00, ...,
       1.28511109e-03, 1.27719252e-03, 2.41553571e-04])

from sklearn.metrics import accuracy_score
accuracy_ls = []
for thres in thresholds:
    y_pred = np.where(final_prediction>thres,1,0)
    accuracy_ls.append(accuracy_score(y_test, y_pred, normalize=True))

accuracy_ls = pd.concat([pd.Series(thresholds),
pd.Series(accuracy_ls)],
                        axis=1)
accuracy_ls.columns = ['thresholds', 'accuracy']
accuracy_ls.sort_values(by='accuracy', ascending=False, inplace=True)
accuracy_ls.head()

   thresholds  accuracy
747    0.498183   0.852300
748    0.498142   0.852207
746    0.498286   0.852207
745    0.499805   0.851835
749    0.496549   0.851835

def plot_roc_curve(fpr, tpr):
    plt.plot(fpr, tpr, color='orange', label='ROC')
    plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--')
    plt.xlabel('False Positive Rate')
    plt.ylabel('True Positive Rate')

```

```
plt.title('Receiver Operating Characteristic (ROC) Curve')  
plt.legend()  
plt.show()  
  
plot_roc_curve(fpr, tpr)
```

