

Household power consumption Regression problem

Problem Statement

- **Regression problem**
 - Collect dataset from here
<https://archive.ics.uci.edu/ml/datasets/individual+household+electric+power+consumption>
 - Here the number of instances is very high, so take a random sample of 50k using the `sample()`.
 - Add all the three columns named `sub_metering_1`, `sub_metering_2` and `sub_metering_3` to get the total energy consumed.
 - Create a Regression model on the basis of attributes.
 - Create Linear Regression, Ridge Regression, Lasso Regression, ElasticNet Regression, Support Vector Regression.
- **Steps to be followed**
 - Data ingestion.
 - EDA (end to end).
 - Preprocessing of the data.
 - Use pickle to store the scaling of the data for later use.
 - Store the final processed data inside MongoDB.
 - Again load the data from MongoDB.
 - Model building.
 - Use GridSearchCV for hyper parameter tuning.
 - Evaluation.
 - R^2 and adjusted R^2 for regression model.

Attribute Information:

1. **date:** Date in format dd/mm/yyyy
2. **time:** time in format hh:mm:ss
3. **global_active_power:** household global minute-averaged active power (in kilowatt)
4. **global_reactive_power:** household global minute-averaged reactive power (in kilowatt)
5. **voltage:** minute-averaged voltage (in volt)
6. **global_intensity:** household global minute-averaged current intensity (in ampere)
7. **sub_metering_1:** energy sub-metering No. 1 (in watt-hour of active energy). It corresponds to the kitchen, containing mainly a dishwasher, an oven and a microwave (hot plates are not electric but gas powered).
8. **sub_metering_2:** energy sub-metering No. 2 (in watt-hour of active energy). It corresponds to the laundry room, containing a washing-machine, a tumble-drier, a refrigerator and a light.
9. **sub_metering_3:** energy sub-metering No. 3 (in watt-hour of active energy). It corresponds to an electric water-heater and an air-conditioner.

1. Data Ingestion:

1.1 Import modules and data to create dataframe # Importing the required libraries

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import warnings
import pymongo

sns.set()
%matplotlib inline
warnings.filterwarnings('ignore')

try:
    read_file = pd.read_csv('dataset/household_power_consumption.txt',
delimiter = ';')
    read_file.to_csv(r"dataset\power_consumption.csv", index=None)
except Exception as err:
    print("Error is: ", err)
else:
    print("File format converted successfully.")
```

File format converted successfully.

1.2 Creating dataframe with random 50000 observations

```
data = pd.read_csv('dataset/power_consumption.csv')
data.shape
```

(2075259, 9)

Note

- Here at the present the number of rows is very high, let's take a sample of 50000 observations

```
df = data.sample(50000)
```

```
df.head()
```

	Date	Time	Global_active_power	Global_reactive_power
511536	6/12/2007	23:00:00	3.118	0.108
1990034	28/9/2010	16:38:00	?	?
550140	2/1/2008	18:24:00	0.628	0.000
1369908	25/7/2009	01:12:00	0.300	0.248

604155	9/2/2008	06:39:00	0.224	0.000
--------	----------	----------	-------	-------

	Voltage	Global_intensity	Sub_metering_1	Sub_metering_2	\
511536	243.110	12.800	0.000	0.000	
1990034	?	?	?	?	
550140	239.030	2.600	0.000	0.000	
1369908	243.610	1.600	0.000	0.000	
604155	243.710	0.800	0.000	0.000	

	Sub_metering_3
511536	18.0
1990034	NaN
550140	0.0
1369908	1.0
604155	0.0

```
df.shape
```

```
(50000, 9)
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 50000 entries, 511536 to 1110791
Data columns (total 9 columns):
```

#	Column	Non-Null Count	Dtype
0	Date	50000 non-null	object
1	Time	50000 non-null	object
2	Global_active_power	50000 non-null	object
3	Global_reactive_power	50000 non-null	object
4	Voltage	50000 non-null	object
5	Global_intensity	50000 non-null	object
6	Sub_metering_1	50000 non-null	object
7	Sub_metering_2	50000 non-null	object
8	Sub_metering_3	49388 non-null	float64

```
dtypes: float64(1), object(8)
```

```
memory usage: 3.8+ MB
```

Observations:

- Now there are 50000 rows and 9 columns (features) in the dataset.
- All the columns except Sub_metering_3 is of object type, even though they have float values.

2. Data Cleaning:

```
# Name of the columns
```

```
df.columns
```

```
Index(['Date', 'Time', 'Global_active_power', 'Global_reactive_power',
      'Voltage', 'Global_intensity', 'Sub_metering_1',
      'Sub_metering_2',
      'Sub_metering_3'],
      dtype='object')
```

2.1 Dropping unnecessary columns

```
df.drop(columns=['Date', 'Time'], axis=1, inplace=True)
df.head()
```

	Global_active_power	Global_reactive_power	Voltage
Global_intensity \			
511536	3.118	0.108	243.110
12.800			
1990034	?	?	?
?			
550140	0.628	0.000	239.030
2.600			
1369908	0.300	0.248	243.610
1.600			
604155	0.224	0.000	243.710
0.800			

	Sub_metering_1	Sub_metering_2	Sub_metering_3
511536	0.000	0.000	18.0
1990034	?	?	NaN
550140	0.000	0.000	0.0
1369908	0.000	0.000	1.0
604155	0.000	0.000	0.0

2.2 Converting data types and replacing special characters

```
for column in df.columns:
    print(f"The unique values in column {column}:")
    print(df[column].unique())
    print(f"\nThe number of unique values in {column} is:
{len(df[column].unique())}")
    print("-----\n")
```

```
The unique values in column Global_active_power:
['1.012' '0.248' '0.308' ... '4.282' '4.316' '5.012']
```

```
The number of unique values in Global_active_power is: 3250
-----
```

```
The unique values in column Global_reactive_power:
['0.190' '0.104' '0.000' '0.164' '0.316' '0.232' '0.046' 0.092 '0.118'
 '0.048' '0.296' '0.292' '0.090' '0.234' '0.324' '0.156' '0.208'
 '0.314'
 '0.132' '0.374' '0.180' '0.236' '0.138' '0.082' '0.084' '0.214'
 '0.198']
```

'0.254' '0.204' '0.226' '0.116' '0.144' '0.148' '0.228' '0.134'
 '0.094'
 '0.400' '0.058' '0.278' '?' '0.068' '0.394' '0.076' '0.072' '0.224'
 '0.274' '0.106' '0.262' '0.092' '0.096' '0.074' '0.078' '0.070'
 '0.240'
 '0.114' '0.080' '0.368' '0.126' '0.124' '0.060' '0.510' '0.416'
 '0.130'
 '0.056' '0.194' '0.098' '0.268' '0.064' '0.454' '0.258' 0.0 '0.052'
 '0.356' '0.220' '0.128' '0.146' '0.062' '0.110' '0.122' '0.230'
 '0.332'
 '0.160' '0.222' '0.050' '0.242' '0.108' '0.174' '0.102' '0.298'
 '0.362'
 '0.066' '0.172' '0.334' '0.320' '0.284' 0.15 '0.088' '0.176' 0.172
 '0.270' '0.154' 0.284 '0.100' '0.168' '0.200' '0.206' '0.322' '0.340'
 '0.210' '0.238' '0.248' '0.366' '0.336' '0.348' '0.054' '0.184'
 '0.412'
 '0.276' '0.432' '0.178' '0.260' '0.196' '0.192' 0.128 '0.212' '0.452'
 '0.436' '0.218' '0.308' '0.272' '0.386' '0.244' '0.112' '0.484'
 '0.382'
 '0.300' 0.242 0.158 '0.376' '0.246' '0.120' '0.286' '0.152' '0.350'
 '0.250' 0.08 '0.344' '0.216' '0.346' '0.256' '0.086' '0.354' '0.188'
 0.064 0.074 '0.266' '0.306' 0.268 '0.202' '0.162' '0.358' '0.142'
 '0.136'
 '0.408' '0.280' '0.538' '0.166' '0.496' 0.114 '0.264' '0.426' '0.372'
 '0.318' '0.150' 0.062 '0.420' '0.330' '0.312' 0.124 '0.478' '0.392'
 '0.378' '0.186' '0.546' '0.370' '0.490' 0.084 0.154 '0.600' '0.338'
 '0.282' '0.380' '0.472' '0.252' 0.06 '0.446' '0.364' '0.908' 0.078
 0.13
 '0.480' '0.182' '0.518' '0.170' '0.140' '0.310' 0.066 '0.462' 0.132
 '0.552' '0.290' 0.116 '0.288' '0.936' '0.390' 0.054 '0.404' '0.516'
 '0.158' '0.428' '0.304' '0.294' '0.402' 0.1 0.258 0.046 '0.342'
 '0.328'
 0.096 '0.388' '0.572' 0.21 '0.476' '0.562' '0.622' '0.574' '0.556'
 0.384
 '0.514' '0.326' '0.360' 0.148 '0.414' 0.098 '0.842' '0.548' 0.294
 '0.302'
 '0.384' '0.540' '0.440' '0.458' '0.470' 0.058 '0.352' '0.444' '0.456'
 '0.398' '0.700' 0.218 '0.450' 0.214 '0.576' '0.488' '0.508' '0.560'
 '0.502' 0.082 0.09 0.146 '0.512' '0.582' '0.748' '0.422' '0.430'
 '0.442'
 '0.424' '0.474' 0.376 '0.770' '0.448' 0.204 '0.564' 0.472 '0.396'
 0.178
 '0.590' 0.056 '0.406' '0.554' 0.684 0.2 0.496 '0.672' 0.152 0.164
 '0.410'
 0.298 0.244 '0.434' 0.094 0.136 '0.530' '0.668' '0.418' 0.184 0.144
 0.206
 0.05 0.22 '0.640' '0.466' 0.226 '0.684' '0.498' '0.438' 0.346 '0.570'
 '0.522' 0.448 '0.708' 0.126 0.138 0.134 0.104 0.068 '0.720' 0.196
 0.166
 0.086 0.108 '0.792' 0.416 '0.670' 0.208 0.122 0.23 '0.604' '0.624'

```

'0.506' 0.052 '0.536' 0.088 '0.652' 0.266 '0.494' 0.102 '0.768' 0.118
0.072 '0.648' '0.460' 0.26 '0.646' 0.278 '0.756' '0.666' 0.222 0.552
'0.802' 0.246 0.112 '0.520' '0.558' 0.07 '0.786' 0.17 0.402 0.168
'0.532'
0.048 0.272 0.12 '0.780' '0.586' '0.762' 0.322 0.182 0.224 0.276
0.076
0.234 0.262 '0.642' '0.486' '0.772' '0.660' '0.644' '0.524' 0.25
'0.592'
'0.482' '0.504' '0.528' '0.500' '0.588' 0.666 '0.464' 0.544 0.14
'0.578'
0.216 '0.610' '1.036' '0.594' '0.534' 0.464 0.312 0.198 0.19 0.194
0.176
'0.568' '0.794' '0.492' 0.65 0.286 0.326 0.458 0.188 '0.664' 0.318
'0.598' 0.238 0.342 0.162 0.202 '0.694' 0.516 0.348 '0.818' '0.550'
'0.698' 0.292 0.636 0.106 '0.732' 0.374 0.18 0.368 0.16 0.212 '0.612'
'0.628' 0.314 0.186 0.236 0.332 0.504 '0.544' 0.192 '0.468' '0.608'
0.142
'0.774' '0.596' '0.742' '0.686' 0.11 0.254 0.394 0.388 '0.618'
'0.584'
0.252 '0.602' '0.734' 0.256 0.522 0.3 0.232 '0.580' '0.626' '0.678'
0.24
'0.778' 0.324 0.39 '0.728' 0.38 '0.632' 0.482 0.27 '0.922' '0.616'
'0.566' 0.288 '0.542' 0.296 0.34 '0.526' 0.174 '0.654' 0.358 0.564
0.33
0.338 '0.730' '0.706' '0.630' '0.606' 0.248 '0.636' '0.656' 0.614
'0.836'
0.37 '0.696' 0.304 0.274 0.398 0.484 0.35 '0.674' '0.614' '0.838'
0.476
'0.998' 0.406 0.282 '0.790' 0.156 '0.902' '0.784' 0.308 '0.718']

```

The number of unique values in Global_reactive_power is: 526

The unique values in column Voltage:

```
['240.290' '241.160' '244.180' ... 235.36 243.14 '251.750']
```

The number of unique values in Voltage is: 2840

The unique values in column Global_intensity:

```
['6.000' '1.000' '1.200' '5.400' '3.200' '2.000' '9.800' '6.400' 1.4
'5.800' '12.000' '7.800' '11.400' '0.800' '8.000' '4.000' '11.000'
'10.800' '4.400' '8.400' '6.800' '2.200' '0.600' '20.000' '1.600'
'6.600'
'11.200' '5.200' '1.400' '4.200' '?' '3.000' '4.600' '5.600' '1.800'
'5.000' '17.400' '6.200' '9.400' '19.600' '9.200' '16.200' '3.800'
'21.000' '10.000' '2.400' '7.000' '7.600' '16.800' 14.0 '15.200'
'7.400'
37.2 '8.200' '3.600' '7.200' '15.800' '11.600' '22.600' '12.600'
'0.200']

```

```

'12.400' 3.4 '2.600' 5.6 2.0 '28.200' '10.400' '9.000' '17.600'
'16.600'
3.0 '3.400' '2.800' '10.200' '18.000' '8.600' '14.600' '19.200'
'9.600'
9.4 '18.400' '20.800' 2.8 '0.400' '8.800' 10.4 '14.800' '16.400'
'14.000'
6.2 '15.600' '10.600' '12.800' 7.2 '19.800' '20.400' '14.400'
'19.400'
'23.000' 6.8 '18.800' 1.6 '21.600' '23.400' '13.000' '27.400' '4.800'
4.0
1.8 '18.200' 8.8 '28.800' '24.800' '11.800' '13.200' '15.400'
'16.000'
5.8 '12.200' 5.0 2.2 '24.400' '21.800' 4.2 '13.400' '13.800' 1.2 7.0
'27.600' 10.0 7.4 '20.200' '13.600' 1.0 '17.800' 10.2 '19.000'
'25.200'
'17.200' '34.200' '24.200' 2.6 13.2 '28.400' '17.000' 11.8 11.4
'31.400'
'14.200' '15.000' 15.2 10.6 8.4 '33.200' '24.600' 6.6 '27.200' 6.0
'24.000' '26.800' 9.2 15.4 '22.400' '34.400' 3.2 12.0 '20.600' 4.4
'18.600' '21.200' 7.6 '29.800' '23.200' 5.4 '22.200' 11.0 '26.200'
15.0
7.8 '25.400' '25.800' 4.8 5.2 '23.800' 2.4 '23.600' 15.6 '31.600'
'29.600' 10.8 '21.400' '26.400' 11.6 13.0 9.6 3.6 '25.600' 6.4
'26.000'
'22.800' 12.8 4.6 '28.000' 43.0 '34.000' '22.000' 8.2 '25.000' 16.0
9.8
13.4 9.0 12.6 '32.600' '32.800' 8.0 '34.600' '29.000' 21.2 16.6
'26.600'
14.4 '33.000' 3.8 14.6 24.0 8.6 '37.600' '30.800' 25.2 11.2 '33.800'
21.4
14.2 '33.600' 12.2 '27.000' 17.8 '28.600' 17.4 27.4 '43.400' '30.400'
'32.400' 15.8 '29.400' '29.200' '30.000' 19.6 12.4 23.4 26.4 24.8
'34.800' '27.800']

```

The number of unique values in Global_intensity is: 256

The unique values in column Sub_metering_1:

```

['11.000' '0.000' 0.0 '1.000' '9.000' '37.000' '2.000' '?' '4.000'
'38.000' 17.0 '36.000' '24.000' '39.000' '35.000' '41.000' 23.0
'23.000'
6.0 '40.000' '45.000' '28.000' '30.000' '13.000' '5.000' '47.000'
'3.000'
'25.000' '43.000' '33.000' '77.000' '16.000' 1.0 '14.000' 37.0
'12.000'
26.0 '10.000' '15.000' '42.000' '22.000' '8.000' '29.000' '55.000'
'7.000' '31.000' 21.0 '32.000' '6.000' '19.000' '74.000' '27.000'
41.0
'49.000' '20.000' '26.000' '56.000' 2.0 13.0 '17.000' '73.000'
'18.000'

```

```

38.0 '44.000' '72.000' '21.000' 43.0 '34.000' '76.000' 24.0 36.0
'48.000'
'78.000' 25.0 '75.000' 35.0 9.0 47.0 '51.000' 15.0 '70.000' 52.0
'53.000'
22.0 '71.000' '67.000' 48.0 '65.000' 32.0 14.0 '46.000' '54.000' 18.0
70.0]

```

The number of unique values in Sub_metering_1 is: 94

The unique values in column Sub_metering_2:

```

['0.000' '2.000' '1.000' 1.0 '4.000' '39.000' '23.000' '35.000' '?'
'36.000' 0.0 '37.000' 68.0 '6.000' '29.000' '15.000' '17.000' 2.0
'5.000'
'21.000' '40.000' '31.000' '28.000' '32.000' '22.000' '38.000'
'3.000'
'20.000' '18.000' '30.000' '11.000' '24.000' '53.000' '76.000'
'43.000'
'34.000' '16.000' '55.000' '27.000' '63.000' '50.000' 17.0 '26.000'
'25.000' '7.000' '73.000' '71.000' '10.000' '19.000' '52.000'
'75.000'
'33.000' '42.000' '68.000' '9.000' '54.000' '14.000' '72.000'
'41.000'
'45.000' '49.000' '66.000' 3.0 '69.000' '12.000' 5.0 '74.000' 36.0
'8.000' 67.0 39.0 35.0 '47.000' '58.000' 4.0 24.0 '13.000' '61.000'
'77.000' '70.000' 71.0 74.0 '62.000' '64.000' 49.0 34.0 '67.000' 38.0
28.0 40.0 '44.000']

```

The number of unique values in Sub_metering_2 is: 91

The unique values in column Sub_metering_3:

```

[ 1.  0. 18. 19.  7. 23. 17. nan 13.  9. 28. 12. 16.  6. 29. 10. 21.
20.
 11.  3.  4. 25.  2.  8. 22. 26. 30.  5. 27. 15. 31. 14. 24.]

```

The number of unique values in Sub_metering_3 is: 33

Observations:

- We have special character ? in columns Sub_metering_1, Sub_metering_2, Global_intensity.
- Also the columns Global_active_power and Voltage have more than 1000 unique values. So we need to check for special characters in them as well.
- We have nan in Sub_metering_3 as well.

To find special characters in these 2 columns

```
df.loc[df['Global_active_power'] == "?", :]
```

	Global_active_power	Global_reactive_power	Voltage
Global_intensity \			
1931603	?	?	?
?			
1618807	?	?	?
?			
1933147	?	?	?
?			
1932297	?	?	?
?			
1936817	?	?	?
?			
...
...			
1989633	?	?	?
?			
1619426	?	?	?
?			
1311597	?	?	?
?			
1930475	?	?	?
?			
1310964	?	?	?
?			

	Sub_metering_1	Sub_metering_2	Sub_metering_3
1931603	?	?	NaN
1618807	?	?	NaN
1933147	?	?	NaN
1932297	?	?	NaN
1936817	?	?	NaN
...
1989633	?	?	NaN
1619426	?	?	NaN
1311597	?	?	NaN
1930475	?	?	NaN
1310964	?	?	NaN

[617 rows x 7 columns]

```
df.loc[df['Voltage'] == "?", :]
```

	Global_active_power	Global_reactive_power	Voltage
Global_intensity \			
1931603	?	?	?
?			

1618807	?	?	?
?			
1933147	?	?	?
?			
1932297	?	?	?
?			
1936817	?	?	?
?			
...
...			
1989633	?	?	?
?			
1619426	?	?	?
?			
1311597	?	?	?
?			
1930475	?	?	?
?			
1310964	?	?	?
?			

	Sub_metering_1	Sub_metering_2	Sub_metering_3
1931603	?	?	NaN
1618807	?	?	NaN
1933147	?	?	NaN
1932297	?	?	NaN
1936817	?	?	NaN
...
1989633	?	?	NaN
1619426	?	?	NaN
1311597	?	?	NaN
1930475	?	?	NaN
1310964	?	?	NaN

[617 rows x 7 columns]

- So yes there are **617** rows where the **?** is present in the dataset
- Also it looks like the sign appears in all the columns at the same time
- As the percentage of these rows is **1%** of the total dataset so we can drop them

Dropping the rows

```
df.drop(df.loc[df['Voltage'] == "?", :].index, inplace=True)
df.shape
```

(49383, 7)

Again checking

```
df.loc[df['Voltage'] == "?", :]
```

```
Empty DataFrame
Columns: [Global_active_power, Global_reactive_power, Voltage,
Global_intensity, Sub_metering_1, Sub_metering_2, Sub_metering_3]
Index: []
```

```
df.loc[df['Global_active_power'] == "?", :]
```

```
Empty DataFrame
Columns: [Global_active_power, Global_reactive_power, Voltage,
Global_intensity, Sub_metering_1, Sub_metering_2, Sub_metering_3]
Index: []
```

Now again checking for nan values

```
df.isnull().sum()
```

```
Global_active_power      0
Global_reactive_power    0
Voltage                  0
Global_intensity          0
Sub_metering_1           0
Sub_metering_2           0
Sub_metering_3           0
dtype: int64
```

Converting the data types

```
df = df.astype({'Global_active_power':float,
'Global_reactive_power':float, 'Voltage':float,
'Global_intensity':float,
'Sub_metering_1':float, 'Sub_metering_2':float})
```

checking the dataset

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 49383 entries, 1928275 to 1332279
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Global_active_power                   49383 non-null  float64
1   Global_reactive_power                 49383 non-null  float64
2   Voltage                               49383 non-null  float64
3   Global_intensity                      49383 non-null  float64
4   Sub_metering_1                       49383 non-null  float64
5   Sub_metering_2                       49383 non-null  float64
6   Sub_metering_3                       49383 non-null  float64
dtypes: float64(7)
memory usage: 3.0 MB
```

2.3 Let's check for duplicate values

```
df[df.duplicated()]
```

	Global_active_power	Global_reactive_power	Voltage	\
950270	0.360	0.078	241.18	
950265	0.360	0.078	241.18	
1047327	0.284	0.048	241.40	
950034	0.360	0.078	241.18	
345658	0.102	0.000	235.86	
...	
845145	0.146	0.000	238.42	
950068	0.360	0.078	241.18	
1914248	0.158	0.000	242.90	
1008734	0.228	0.000	244.39	
872385	0.082	0.000	242.34	

	Global_intensity	Sub_metering_1	Sub_metering_2	Sub_metering_3
950270	1.6	0.0	0.0	1.0
950265	1.6	0.0	0.0	1.0
1047327	1.2	0.0	0.0	0.0
950034	1.6	0.0	0.0	1.0
345658	0.4	0.0	0.0	0.0
...
845145	0.6	0.0	0.0	1.0
950068	1.6	0.0	0.0	1.0
1914248	0.6	0.0	0.0	0.0
1008734	1.0	0.0	0.0	0.0
872385	0.2	0.0	0.0	1.0

```
[194 rows x 7 columns]
```

```
# Dropping the duplicated values as well
```

```
df.drop_duplicates(inplace=True)
```

```
# Creating a new column for 'total energy consumed'
```

```
# Then removing the columns 'Sub_metering_1', 'Sub_metering_2' and  
'Sub_metering_3'
```

```
df["Total_energy_consumed"] = df['Sub_metering_1'] +
df['Sub_metering_2'] + df['Sub_metering_3']
df.drop(columns=['Sub_metering_1', 'Sub_metering_2',
'Sub_metering_3'], axis=1, inplace=True)
df.head()
```

	Global_active_power	Global_reactive_power	Voltage	\
1928275	1.012	0.190	240.29	
1892037	0.248	0.104	241.16	
533121	0.308	0.000	244.18	
725441	1.314	0.000	240.09	
1078558	0.754	0.164	243.22	

	Global_intensity	Total_energy_consumed
1928275	6.0	12.0
1892037	1.0	2.0
533121	1.2	0.0
725441	5.4	18.0
1078558	3.2	0.0

```
df.shape
```

```
(49189, 5)
```

3. Exploratory data analysis

3.1 Basic Profile of the data

Checking the details of the dataframe

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 49189 entries, 1928275 to 1332279
Data columns (total 5 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Global_active_power                   49189 non-null  float64
1   Global_reactive_power                 49189 non-null  float64
2   Voltage                               49189 non-null  float64
3   Global_intensity                      49189 non-null  float64
4   Total_energy_consumed                 49189 non-null  float64
dtypes: float64(5)
memory usage: 2.3 MB
```

Differentiating numerical and categorical columns

```
numerical_features = [feature for feature in df.columns if
df[feature].dtypes != 'O']
categorical_features = [feature for feature in df.columns if
df[feature].dtypes == 'O']
```

```
print(f"The number of Numerical features are:
{len(numerical_features)}, and the column names are:\n{numerical_features}")
print(f"\nThe number of Categorical features are:
{len(categorical_features)}, and the column names are:\n{categorical_features}")
```

The number of Numerical features are: 5, and the column names are:
['Global_active_power', 'Global_reactive_power', 'Voltage',
'Global_intensity', 'Total_energy_consumed']

The number of Categorical features are: 0, and the column names are:
[]

Observations:

- Now we have 49189 rows with no null and duplicate values and all the 5 columns have numerical (float) data type.

3.2 Statistical Analysis of the data

```
df.describe().T
```

	count	mean	std	min
25% \				
Global_active_power	49189.0	1.093096	1.056170	0.078
0.308				
Global_reactive_power	49189.0	0.124045	0.112738	0.000
0.048				
Voltage	49189.0	240.839414	3.246788	225.720
238.980				
Global_intensity	49189.0	4.634906	4.438016	0.200
1.400				
Total_energy_consumed	49189.0	8.937466	12.955325	0.000
0.000				

	50%	75%	max
Global_active_power	0.606	1.528	9.938
Global_reactive_power	0.100	0.194	1.036
Voltage	241.030	242.900	253.100
Global_intensity	2.600	6.400	43.400
Total_energy_consumed	1.000	18.000	126.000

Observations:

- There are possible Outliers in columns Global_active_power, Global_intensity, Total_energy_consumed.

3.3 Graphical Analysis of the data


3.3.1 Univariate Analysis

For numerical features
Kernal Density plots

```
plt.figure(figsize=(15, 15))  
plt.suptitle('Univariate Analysis of Numerical Features', fontsize=20,  
fontweight='bold', alpha=0.8, y=1.)
```

```
for i in range(0, len(numerical_features)):  
    plt.subplot(3, 2, i+1)  
    sns.kdeplot(x=df[numerical_features[i]], shade=True, color='r')  
    plt.xlabel(numerical_features[i], fontsize=15)  
    plt.tight_layout()
```

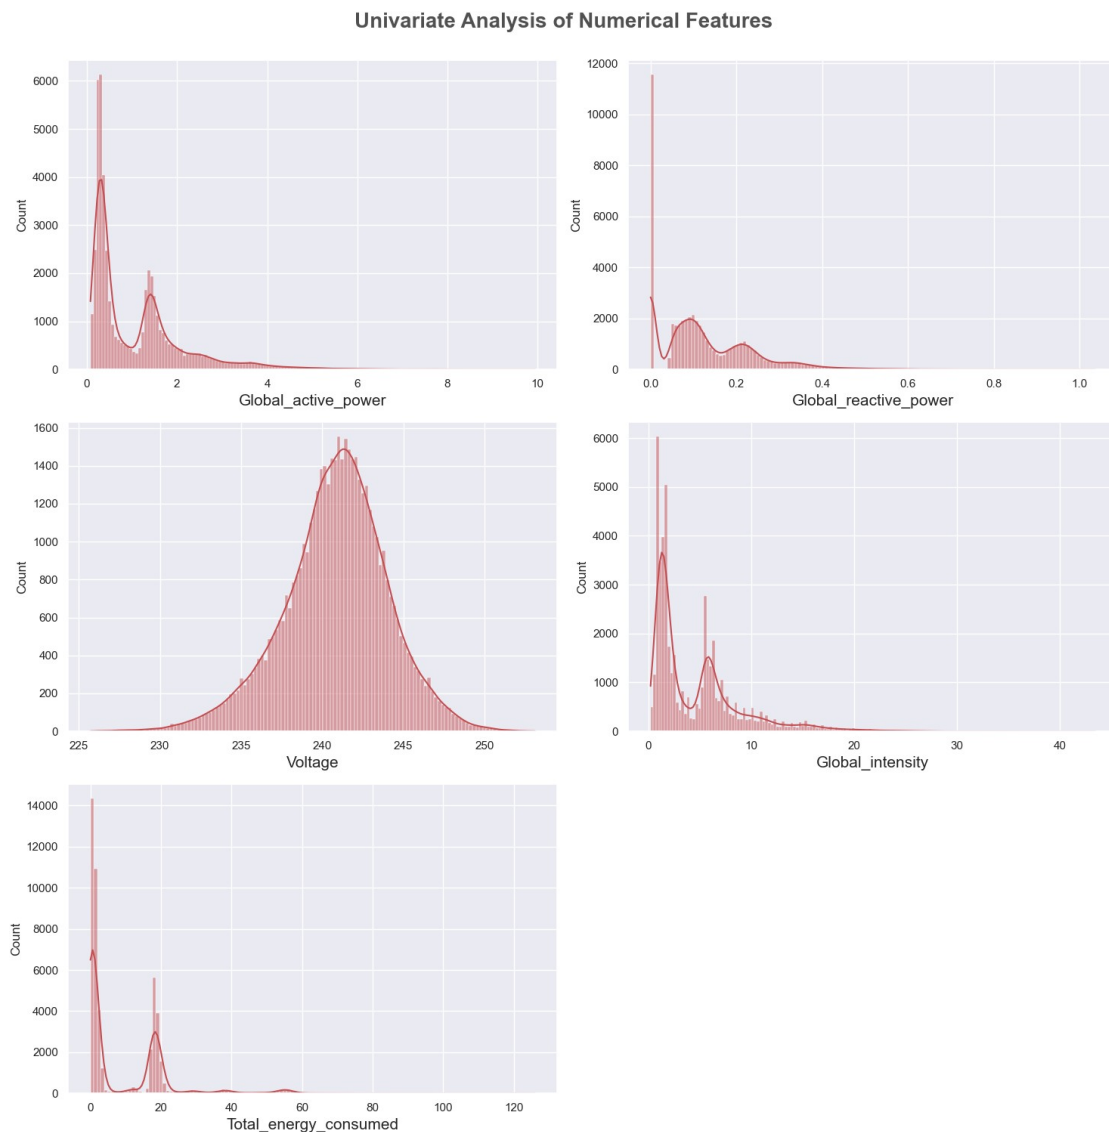
Univariate Analysis of Numerical Features



Histograms

```
plt.figure(figsize=(15, 15))
plt.suptitle('Univariate Analysis of Numerical Features', fontsize=20,
fontweight='bold', alpha=0.8, y=1.)
```

```
for i in range(0, len(numerical_features)):
    plt.subplot(3, 2, i+1)
    sns.histplot(x=df[numerical_features[i]], kde=True, color='r')
    plt.xlabel(numerical_features[i], fontsize=15)
    plt.tight_layout()
```



Observations:

- Only Voltage has normal distribution.
- All other columns are right skewed and they may have outliers.

- Too many values near to 0 in Global_active_power, Global_reactive_power, Global_intensity and Total_energy_consumed columns.

3.3.2 Multivariate Analysis

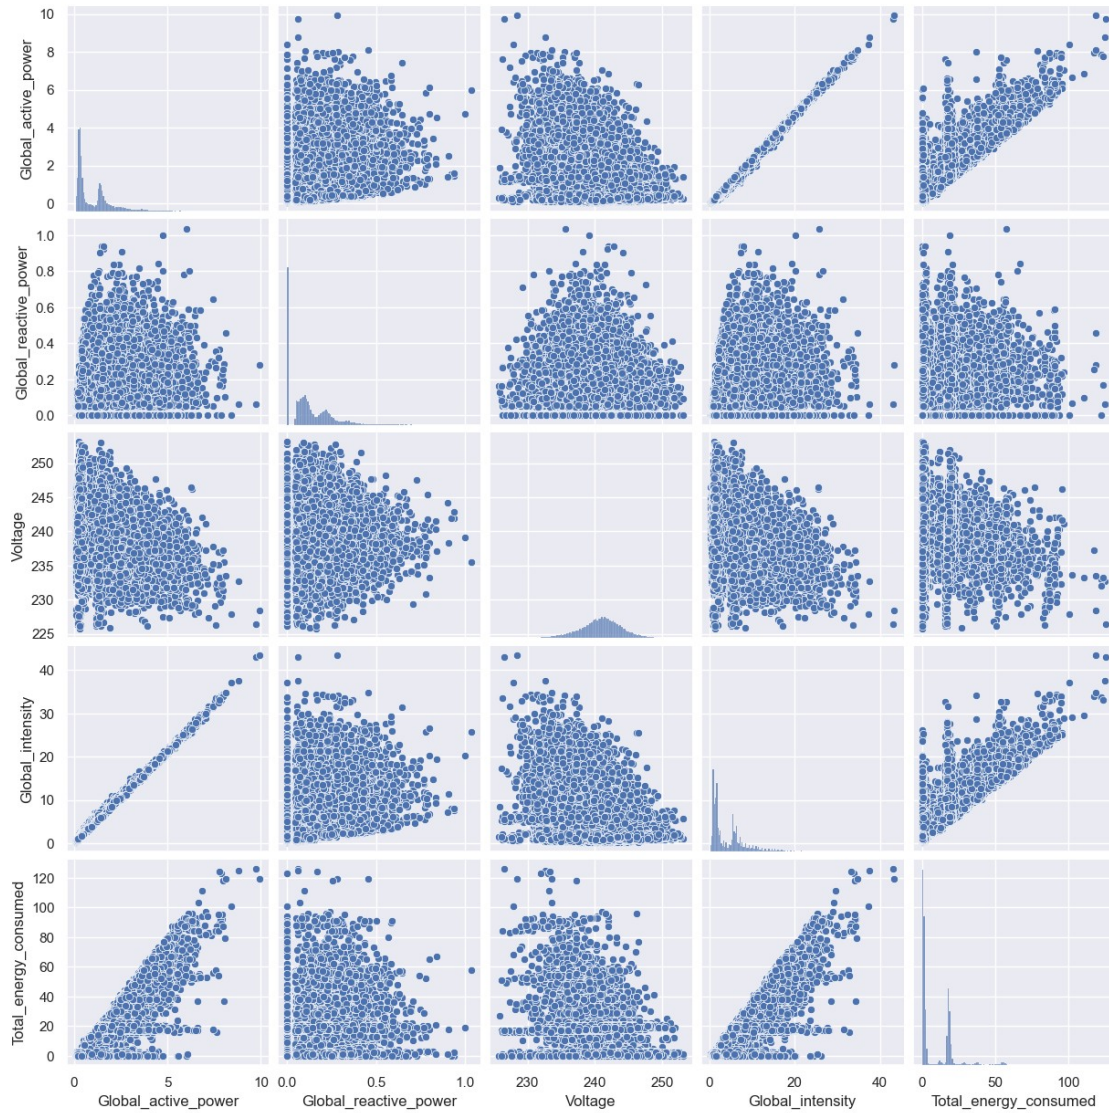
Checking Multicollinearity in the numerical features

```
df[list(df[numerical_features].columns)].corr()
```


	Global_active_power	Global_reactive_power
Voltage \		
Global_active_power	1.000000	0.231006 -
0.397982		
Global_reactive_power	0.231006	1.000000 -
0.103893		
Voltage	-0.397982	-0.103893
1.000000		
Global_intensity	0.998883	0.249991 -
0.409457		
Total_energy_consumed	0.848864	0.166474 -
0.342788		
	Global_intensity	Total_energy_consumed
Global_active_power	0.998883	0.848864
Global_reactive_power	0.249991	0.166474
Voltage	-0.409457	-0.342788
Global_intensity	1.000000	0.845797
Total_energy_consumed	0.845797	1.000000

Graphical representation

```
sns.pairplot(df[numerical_features])
plt.show()
```



```
sns.set(rc={'figure.figsize':(15,10)})
sns.heatmap(df[numerical_features].corr(), cmap='CMRmap', annot=True)
plt.show()
```



Observations:

- Global_intensity and Global_active_power is completely correlated.
- Total_energy_consumed is also highly correlated with Global_intensity and Global_active_power.

4. Data Pre-Processing

4.1 Number of unique values in each column

```
df.nunique()
```

```
Global_active_power    2659
Global_reactive_power    352
Voltage                2097
Global_intensity        170
Total_energy_consumed   107
dtype: int64
```

4.2 Outlier handling

Creating a function to detect outliers

```
def detect_outliers(col):
    percentile25 = df[col].quantile(0.25)
    percentile75 = df[col].quantile(0.75)
    print('\n ####', col , '####')
```

```

print("25percentile: ",percentile25)
print("75percentile: ",percentile75)
iqr = percentile75 - percentile25
upper_limit = percentile75 + 1.5 * iqr
lower_limit = percentile25 - 1.5 * iqr
print("Upper limit: ",upper_limit)
print("Lower limit: ",lower_limit)
df.loc[(df[col]>upper_limit), col]= upper_limit
df.loc[(df[col]<lower_limit), col]= lower_limit
return df

```

Now applying the function on all the columns as all are of continupus type

```

for col in numerical_features:
    detect_outliers(col)

```

```

#### Global_active_power ####
25percentile: 0.308
75percentile: 1.528
Upper limit: 3.358
Lower limit: -1.522

```

```

#### Global_reactive_power ####
25percentile: 0.048
75percentile: 0.194
Upper limit: 0.41300000000000003
Lower limit: -0.17100000000000004

```

```

#### Voltage ####
25percentile: 238.98
75percentile: 242.9
Upper limit: 248.78000000000003
Lower limit: 233.09999999999997

```

```

#### Global_intensity ####
25percentile: 1.4
75percentile: 6.4
Upper limit: 13.9
Lower limit: -6.1

```

```

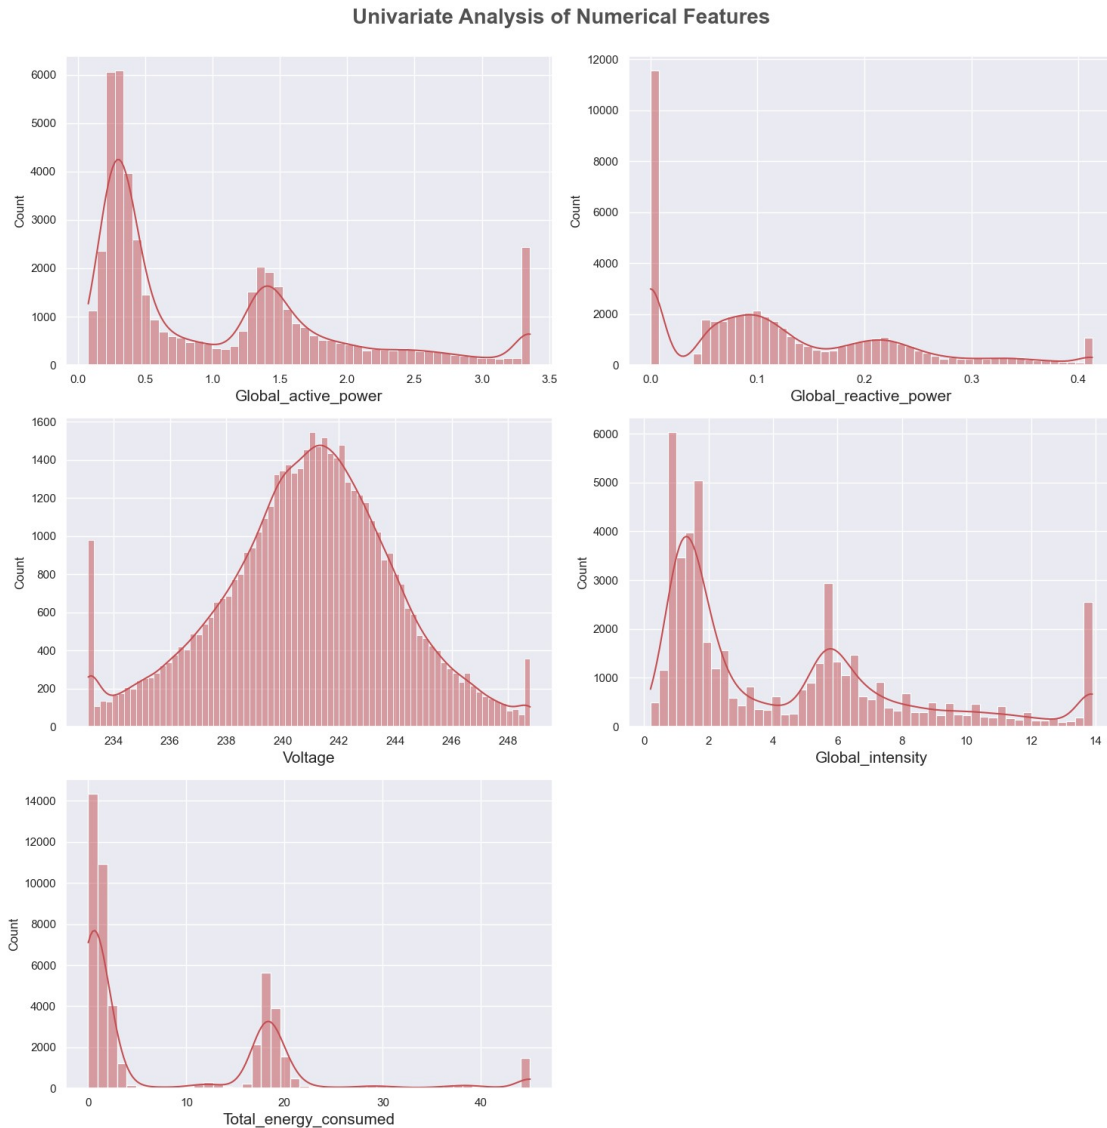
#### Total_energy_consumed ####
25percentile: 0.0
75percentile: 18.0
Upper limit: 45.0
Lower limit: -27.0

```

Again checking

```
plt.figure(figsize=(15, 15))
plt.suptitle('Univariate Analysis of Numerical Features', fontsize=20,
fontweight='bold', alpha=0.8, y=1.)
```

```
for i in range(0, len(numerical_features)):
    plt.subplot(3, 2, i+1)
    sns.histplot(x=df[numerical_features[i]], kde=True, color='r')
    plt.xlabel(numerical_features[i], fontsize=15)
    plt.tight_layout()
```



```
fig, ax = plt.subplots(figsize=(15,10))
plt.suptitle('Finding Outliers in Numerical Features', fontsize=20,
fontweight='bold', alpha=0.8, y=1.)
sns.boxplot(data=df[numerical_features], width= 0.5, ax=ax,
fliersize=3)
plt.show()
```

Finding Outliers in Numerical Features



Observations:

- Now we can see that the outliers are gone.

Let's save the clean data to a folder and then to mongodb for later use

```
try:
    df.to_csv("dataset/power_consumption_cleaned.csv", index=None)
except Exception as err:
    print("Error is: ", err)
else:
    print("Clean csv file created successfully.")
```

Clean csv file created successfully.

converting to json file

```
df2 = pd.read_csv('dataset/power_consumption_cleaned.csv')
```

```
try:
    df2.to_json('dataset/power_consumption_cleaned.json')
except Exception as err:
    print("Error is: ", err)
else:
    print("Json file created successfully.")
```

Json file created successfully.

MongoDB part

Checking the file

```
df_json = pd.read_json('dataset/power_consumption_cleaned.json')
df_json.head()
```

	Global_active_power	Global_reactive_power	Voltage
Global_intensity \			
0	1.012	0.190	240.29
6.0			
1	0.248	0.104	241.16
1.0			
2	0.308	0.000	244.18
1.2			
3	1.314	0.000	240.09
5.4			
4	0.754	0.164	243.22
3.2			

	Total_energy_consumed
0	12
1	2
2	0
3	18
4	0

connecting with the server

```
try:
    client =
pymongo.MongoClient("mongodb+srv://ineuron:Project1@cluster0.rp4qzrr.m
ongodb.net/?retryWrites=true&w=majority")
except Exception as e:
    print(e)
else:
    print("Connection to MongoDB server is successful.")
```

Connection to MongoDB server is successful.

creating database and collection

```
db = client["household_power_consumption"]
coll = db['power_consumption']
```

```
try:
    import json
except ImportError:
    import simplejson as json
```

```
# Inserting the data into the collection
```

```
try:
    with open('dataset/power_consumption_cleaned.json') as file:
        file_data = json.load(file)

        coll.insert_many([file_data])
except Exception as e:
    print(e)
else:
    print("Data inserted successfully.")
```

Data inserted successfully.

```
Loading the data from MongoDB
```

```
# Now to read the data
```

```
# importing the library to take care of the objectid created by  
mongodb
```

```
import bson.json_util as json_util
```

```
results = coll.find()
```

```
try:
    for result in results:
        data = json_util.dumps(result)
        clean_df = pd.read_json(data, orient='index')
except Exception as e:
    print(e)
else:
    clean_df
```

```
clean_df
```

	\$oid	0	1
2 \			
_id	63628dbb4240b6a33e820c11	NaN	NaN
NaN			
Global_active_power	NaN	1.012	0.248
0.308			
Global_reactive_power	NaN	0.190	0.104
0.000			
Voltage	NaN	240.290	241.160
244.180			
Global_intensity	NaN	6.000	1.000
1.200			
Total_energy_consumed	NaN	12.000	2.000
0.000			
	3	4	5
8 \			
	6	7	

_id	NaN	NaN	NaN	NaN	NaN
Global_active_power	1.314	0.754	0.22	0.378	2.264
Global_reactive_power	0.000	0.164	0.00	0.316	0.232
Voltage	240.090	243.220	240.16	237.630	238.660
Global_intensity	5.400	3.200	1.00	2.000	9.800
Total_energy_consumed	18.000	0.000	0.00	1.000	20.000

\	...	49179	49180	49181	49182	49183
_id	...	NaN	NaN	NaN	NaN	NaN
Global_active_power	...	2.256	0.414	0.30	0.188	1.422
Global_reactive_power	...	0.104	0.096	0.11	0.050	0.058
Voltage	...	239.810	240.420	243.35	240.080	237.520
Global_intensity	...	9.800	1.800	1.40	0.800	6.000
Total_energy_consumed	...	30.000	0.000	1.00	1.000	17.000

_id	49184	49185	49186	49187	49188
Global_active_power	1.260	0.196	1.494	0.376	1.404
Global_reactive_power	0.408	0.128	0.144	0.000	0.144
Voltage	243.280	237.060	241.470	244.820	238.840
Global_intensity	5.400	1.000	6.200	1.600	5.800
Total_energy_consumed	14.000	0.000	19.000	0.000	1.000

[6 rows x 49190 columns]

```
clean_df = clean_df.transpose()
clean_df.head()
```

	_id	Global_active_power
Global_reactive_power \		
\$oid 63628dbb4240b6a33e820c11		NaN
NaN		
0	NaN	1.012
0.19		
1	NaN	0.248
0.104		
2	NaN	0.308

0.0		
3	NaN	1.314
0.0		

	Voltage	Global_intensity	Total_energy_consumed
\$oid	NaN	NaN	NaN
0	240.29	6.0	12.0
1	241.16	1.0	2.0
2	244.18	1.2	0.0
3	240.09	5.4	18.0

Removing the column `_id` and row `oid` as we donot need them

Removing the '_id' column

```
clean_df.drop(['_id'], axis=1, inplace=True)
clean_df
```

	Global_active_power	Global_reactive_power	Voltage
Global_intensity \			
\$oid	NaN	NaN	NaN
NaN			
0	1.012	0.19	240.29
6.0			
1	0.248	0.104	241.16
1.0			
2	0.308	0.0	244.18
1.2			
3	1.314	0.0	240.09
5.4			
...
..			
49184	1.26	0.408	243.28
5.4			
49185	0.196	0.128	237.06
1.0			
49186	1.494	0.144	241.47
6.2			
49187	0.376	0.0	244.82
1.6			
49188	1.404	0.144	238.84
5.8			

	Total_energy_consumed
\$oid	NaN
0	12.0
1	2.0
2	0.0
3	18.0
...	...

```

49184          14.0
49185           0.0
49186          19.0
49187           0.0
49188           1.0

```

```
[49190 rows x 5 columns]
```

```
# Again transposing so we can get the '$oid' as a column
```

```

clean_df = clean_df.transpose()
clean_df.head()

```

	\$oid	0	1	2	3	4
5 \						
Global_active_power	NaN	1.012	0.248	0.308	1.314	0.754
0.22						
Global_reactive_power	NaN	0.19	0.104	0.0	0.0	0.164
0.0						
Voltage	NaN	240.29	241.16	244.18	240.09	243.22
240.16						
Global_intensity	NaN	6.0	1.0	1.2	5.4	3.2
1.0						
Total_energy_consumed	NaN	12.0	2.0	0.0	18.0	0.0
0.0						
	6	7	8	...	49179	49180
49181 \						
Global_active_power	0.378	2.264	1.52	...	2.256	0.414
0.3						
Global_reactive_power	0.316	0.232	0.046	...	0.104	0.096
0.11						
Voltage	237.63	238.66	239.32	...	239.81	240.42
243.35						
Global_intensity	2.0	9.8	6.4	...	9.8	1.8
1.4						
Total_energy_consumed	1.0	20.0	18.0	...	30.0	0.0
1.0						
	49182	49183	49184	49185	49186	49187
49188						
Global_active_power	0.188	1.422	1.26	0.196	1.494	0.376
1.404						
Global_reactive_power	0.05	0.058	0.408	0.128	0.144	0.0
0.144						
Voltage	240.08	237.52	243.28	237.06	241.47	244.82
238.84						
Global_intensity	0.8	6.0	5.4	1.0	6.2	1.6
5.8						
Total_energy_consumed	1.0	17.0	14.0	0.0	19.0	0.0

1.0

[5 rows x 49190 columns]

Removing the '\$oid' column

```
clean_df.drop(['$oid'], axis=1, inplace=True)
clean_df
```

	0	1	2	3	4	5
6 \						
Global_active_power	1.012	0.248	0.308	1.314	0.754	0.22
0.378						
Global_reactive_power	0.19	0.104	0.0	0.0	0.164	0.0
0.316						
Voltage	240.29	241.16	244.18	240.09	243.22	240.16
237.63						
Global_intensity	6.0	1.0	1.2	5.4	3.2	1.0
2.0						
Total_energy_consumed	12.0	2.0	0.0	18.0	0.0	0.0
1.0						

	7	8	9	...	49179	49180
49181 \						
Global_active_power	2.264	1.52	0.342	...	2.256	0.414
0.3						
Global_reactive_power	0.232	0.046	0.092	...	0.104	0.096
0.11						
Voltage	238.66	239.32	244.61	...	239.81	240.42
243.35						
Global_intensity	9.8	6.4	1.4	...	9.8	1.8
1.4						
Total_energy_consumed	20.0	18.0	1.0	...	30.0	0.0
1.0						

	49182	49183	49184	49185	49186	49187
49188						
Global_active_power	0.188	1.422	1.26	0.196	1.494	0.376
1.404						
Global_reactive_power	0.05	0.058	0.408	0.128	0.144	0.0
0.144						
Voltage	240.08	237.52	243.28	237.06	241.47	244.82
238.84						
Global_intensity	0.8	6.0	5.4	1.0	6.2	1.6
5.8						
Total_energy_consumed	1.0	17.0	14.0	0.0	19.0	0.0
1.0						

[5 rows x 49189 columns]

```
# getting the actual dataframe
```

```
clean_df = clean_df.transpose()
clean_df.head()
```

```
Global_active_power Global_reactive_power Voltage
Global_intensity \
0          1.012          0.19  240.29          6.0
1          0.248          0.104  241.16          1.0
2          0.308          0.0  244.18          1.2
3          1.314          0.0  240.09          5.4
4          0.754          0.164  243.22          3.2
```

```
Total_energy_consumed
0          12.0
1           2.0
2           0.0
3          18.0
4           0.0
```

```
final_df = clean_df.copy()
final_df
```

```
Global_active_power Global_reactive_power Voltage
Global_intensity \
0          1.012          0.19  240.29
6.0
1          0.248          0.104  241.16
1.0
2          0.308          0.0  244.18
1.2
3          1.314          0.0  240.09
5.4
4          0.754          0.164  243.22
3.2
...          ...          ...          .
..
49184          1.26          0.408  243.28
5.4
49185          0.196          0.128  237.06
1.0
49186          1.494          0.144  241.47
6.2
49187          0.376          0.0  244.82
1.6
```

49188	1.404	0.144	238.84
5.8			

	Total_energy_consumed
0	12.0
1	2.0
2	0.0
3	18.0
4	0.0
...	...
49184	14.0
49185	0.0
49186	19.0
49187	0.0
49188	1.0

[49189 rows x 5 columns]

4.3 Creating independent and dependent variables

Split X and y

- **Split Dataframe to X and y**
- **Here we set a variable X i.e, independent columns, and a variable y i.e, dependent column as the Total_energy_consumed column.**

```
X = final_df.drop("Total_energy_consumed", axis=1)
```

```
y = final_df["Total_energy_consumed"]
```

```
# Checking the independent and dependent variables
```

```
X.head()
```

	Global_active_power	Global_reactive_power	Voltage	Global_intensity
0	1.012	0.19	240.29	6.0
1	0.248	0.104	241.16	1.0
2	0.308	0.0	244.18	1.2
3	1.314	0.0	240.09	5.4
4	0.754	0.164	243.22	3.2

```
y.head()
```

0	12.0
1	2.0
2	0.0
3	18.0
4	0.0

```
Name: Total_energy_consumed, dtype: object
```

```
# Doing Test Train split
```

```
from sklearn.model_selection import train_test_split
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.33, random_state=42)
```

Let's see the datasets

```
X_train.head()
```

	Global_active_power	Global_reactive_power	Voltage
Global_intensity			
8331	3.358	0.08	238.96
13.9			
12276	0.184	0.0	242.04
0.8			
18649	0.33	0.096	239.5
1.4			
4127	0.2	0.078	242.01
1.0			
20274	1.246	0.068	241.46
5.2			

```
y_train.head()
```

8331	45.0
12276	0.0
18649	2.0
4127	0.0
20274	1.0

Name: Total_energy_consumed, dtype: object

```
X_test.head()
```

	Global_active_power	Global_reactive_power	Voltage
Global_intensity			
38057	0.308	0.084	243.6
1.2			
43453	1.78	0.132	236.75
7.4			
27271	2.044	0.0	243.94
8.4			
45630	0.452	0.074	241.73
2.0			
28595	0.226	0.142	242.24
1.0			

```
y_test.head()
```

38057	2.0
43453	18.0
27271	20.0
45630	0.0
28595	0.0

Name: Total_energy_consumed, dtype: object

Let's check the shapes of each datasets

```
X_train.shape
```

```
(32956, 4)
```

```
X_test.shape
```

```
(16233, 4)
```

Observations:

- So now we have 32956 rows for training and 16233 for test datasets.

4.4 Standardizing or feature scaling the dataset

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()  
scaler
```

```
StandardScaler()
```

```
# calculate the mean and std dev
```

```
# Here we are fitting only the training data without transforming
```

```
scale = scaler.fit(X_train)  
scale
```

```
StandardScaler()
```

```
# Printing the mean
```

```
print(scale.mean_)
```

```
[1.05394908e+00 1.22367004e-01 2.40849102e+02 4.45267326e+00]
```

Saving the scale to use it later to transform the data and predict the values

```
# To save a Standard scaler object
```

```
import pickle
```

```
with open('scaled.pkl', 'wb') as f:  
    pickle.dump(scale, f)
```

```
# Loading the scaled object to transform the data
```

```
with open('scaled.pkl', 'rb') as f:  
    scaled = pickle.load(f)
```

```
# Now transforming the train and test dataset
```

```
X_train_tf = scaled.transform(X_train)  
X_test_tf = scaled.transform(X_test)
```



```
# Checking the transformed data
```

```
X_train_tf
```

```
array([[ 2.51411749, -0.39834171, -0.60179369,  2.48867721],
       [-0.9492647 , -1.15051518,  0.3793731 , -0.96221132],
       [-0.78995349, -0.24790702, -0.42977094, -0.80415535],
       ...,
       [-0.30983749,  0.25981008,  2.52647185, -0.38267279],
       [ 0.24884294, -1.15051518, -0.71010431,  0.24955107],
       [-0.80741225, -0.04105931,  0.02258517, -0.80415535]])
```

```
X_test_tf
```

```
array([[ -0.81395928, -0.36073304,  0.8763277 , -0.85684068],
       [ 0.79224695,  0.09057105, -1.30581272,  0.77640428],
       [ 1.08031654, -1.15051518,  0.98463832,  1.03983088],
       ...,
       [-0.9492647 , -1.15051518,  0.15000943, -0.96221132],
       [-0.85105916, -1.15051518,  0.93048301, -0.85684068],
       [-0.74412423,  0.86154885, -0.2577482 , -0.69878471]])
```

5. Model Building

5.1 Import required packages for model training

```
from sklearn.linear_model import LinearRegression, Ridge, Lasso,
ElasticNet
from sklearn.svm import SVR
from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score
```

5.2 Create a Function to evaluate all the models

```
def evaluate_model(true, predicted, X_test_tf):
    mae = mean_absolute_error(true, predicted)
    mse = mean_squared_error(true, predicted)
    rmse = np.sqrt(mean_squared_error(true, predicted))
    r2_square = r2_score(true, predicted)
    adj_r2 = 1 - (1 - r2_square)*(len(true)-1)/(len(true) -
X_test_tf.shape[1] - 1)
    return mae, rmse, r2_square, adj_r2
```

```
models = {
    "Linear Regression": LinearRegression(),
    "Lasso": Lasso(),
    "Ridge": Ridge(),
    "Elastic": ElasticNet(),
    "svr": SVR()
}
```

```
model_list = []
r2_list = []
```

```

adj_r2_list = []

for i in range(len(list(models))):
    model = list(models.values())[i]

    # Train model
    model.fit(X_train_tf, y_train)

    # Make predictions
    y_train_pred = model.predict(X_train_tf)
    y_test_pred = model.predict(X_test_tf)

    # Evaluate Train and Test dataset
    model_train_mae, model_train_rmse, model_train_r2,
    model_train_adjusted_r2 = evaluate_model(y_train, y_train_pred,
    X_test_tf)

    model_test_mae, model_test_rmse, model_test_r2,
    model_test_adjusted_r2 = evaluate_model(y_test, y_test_pred,
    X_test_tf)

    print(list(models.keys())[i])
    model_list.append(list(models.keys())[i])

    print('Model performance for Training set')
    print("- Root Mean Squared Error:
{:.4f}".format(model_train_rmse))
    print("- Mean Absolute Error: {:.4f}".format(model_train_mae))
    print("- R2 Score: {:.4f}".format(model_train_r2))
    print("- Adjusted R2 Score:
{:.4f}".format(model_train_adjusted_r2))

    print('-----')

    print('Model performance for Test set')
    print("- Root Mean Squared Error: {:.4f}".format(model_test_rmse))
    print("- Mean Absolute Error: {:.4f}".format(model_test_mae))
    print("- R2 Score: {:.4f}".format(model_test_r2))
    print("- Adjusted R2 Score:
{:.4f}".format(model_test_adjusted_r2))
    r2_list.append(model_test_r2)
    adj_r2_list.append(model_test_adjusted_r2)

    print('='*50)
    print('\n')

```

Linear Regression

Model performance for Training set

- Root Mean Squared Error: 6.2380

- Mean Absolute Error: 4.1482
- R2 Score: 0.6999
- Adjusted R2 Score: 0.6999

Model performance for Test set

- Root Mean Squared Error: 6.1659
- Mean Absolute Error: 4.0543
- R2 Score: 0.7061
- Adjusted R2 Score: 0.7060

=====

Lasso

Model performance for Training set

- Root Mean Squared Error: 6.3860
- Mean Absolute Error: 4.4077
- R2 Score: 0.6855
- Adjusted R2 Score: 0.6855

Model performance for Test set

- Root Mean Squared Error: 6.3039
- Mean Absolute Error: 4.3113
- R2 Score: 0.6928
- Adjusted R2 Score: 0.6927

=====

Ridge

Model performance for Training set

- Root Mean Squared Error: 6.2381
- Mean Absolute Error: 4.1472
- R2 Score: 0.6999
- Adjusted R2 Score: 0.6999

Model performance for Test set

- Root Mean Squared Error: 6.1656
- Mean Absolute Error: 4.0530
- R2 Score: 0.7061
- Adjusted R2 Score: 0.7060

=====

Elastic

Model performance for Training set

- Root Mean Squared Error: 6.7325
- Mean Absolute Error: 4.9730
- R2 Score: 0.6504
- Adjusted R2 Score: 0.6504

Model performance for Test set

- Root Mean Squared Error: 6.6634
- Mean Absolute Error: 4.8884
- R2 Score: 0.6567
- Adjusted R2 Score: 0.6567

```
svr
Model performance for Training set
- Root Mean Squared Error: 6.1315
- Mean Absolute Error: 3.0420
- R2 Score: 0.7101
- Adjusted R2 Score: 0.7100
```

```
-----
Model performance for Test set
- Root Mean Squared Error: 6.1046
- Mean Absolute Error: 2.9965
- R2 Score: 0.7119
- Adjusted R2 Score: 0.7118
=====
```

5.3 Results of all Models

```
adj_r2_list
```

```
[0.7059998203542576,
 0.6926925692846377,
 0.7060327002780958,
 0.6566502290168923,
 0.7118220744716512]
```

```
pd.DataFrame(list(zip(model_list, r2_list, adj_r2_list)),
              columns=['Model Name', 'R2_Score', 'Adjusted
R2_Score']).sort_values(by=["R2_Score"], ascending=False)
```

	Model Name	R2_Score	Adjusted R2_Score
4	svr	0.711893	0.711822
2	Ridge	0.706105	0.706033
0	Linear Regression	0.706072	0.706000
1	Lasso	0.692768	0.692693
3	Elastic	0.656735	0.656650

Observations:

- We can see the best Adjusted R2 value is of the SVR model but the SVR, Ridge, Linear Regression models are also very close.
- So now we can use SVR for Hyper Parameter Tuning to find it's best values.

5.4 Hyper Parameter Tuning (using GridSearchCV)

importing the library

```
from sklearn.model_selection import GridSearchCV
```

Creating the svr model

```
svr = SVR()
```

```
svr
```

```
SVR()
```

training the model

```
svr.fit(X_train_tf, y_train)
```

```
SVR()
```

```
params = {'kernel': ('linear', 'rbf')}
```

```
grid = GridSearchCV(estimator=svr, param_grid=params, cv=3, verbose=2,  
n_jobs=-1)
```

```
grid.fit(X_train_tf, y_train)
```

```
print(grid.best_params_)
```

Fitting 3 folds for each of 2 candidates, totalling 6 fits
{'kernel': 'rbf'}

5.5 Training the model with best Parameters

```
best_model = SVR(kernel='rbf')
```

```
best_model
```

```
SVR()
```

training the model

```
best_model.fit(X_train_tf, y_train)
```

```
SVR()
```

5.6 Saving the optimized model for later usage

saving the model

```
with open('model.pkl', 'wb') as f:  
    pickle.dump(best_model, f)
```

Testing the model with new data to get prediction

Inserting data from outside

```
Global_active_power = float(input("Enter The value: "))
```

```
Global_reactive_power = float(input("Enter The value: "))
```

```

Voltage = float(input("Enter The value: "))
Global_intensity = float(input("Enter The value: "))

test_set = {'Global_active_power': Global_active_power,
            'Global_reactive_power': Global_reactive_power,
            'Voltage': Voltage, 'Global_intensity': Global_intensity}
print("\nThe entered values are:\n")
print(test_set)

```

```

Enter The value: 0.300
Enter The value: 0.248
Enter The value: 243.610
Enter The value: 1.600

```

The entered values are:

```

{'Global_active_power': 0.3, 'Global_reactive_power': 0.248,
'Voltage': 243.61, 'Global_intensity': 1.6}

```

creating a dataframe of the entered data

```

test_set = pd.DataFrame(test_set, index=[1])
test_set

```

```

      Global_active_power  Global_reactive_power  Voltage
Global_intensity
1                0.3                0.248    243.61
1.6

```

```
test_set.shape
```

```
(1, 4)
```

Loading the scaled object to transform the entered data

```

with open('scaled.pkl', 'rb') as f:
    scaled = pickle.load(f)

```

Now transforming the entered data

```

test_set_tf = scaled.transform(test_set)
test_set_tf

```

```
array([[ -0.82268867,  1.18122257,  0.87951331, -0.75147003]])
```

loading the model

```

with open('model.pkl', 'rb') as f:
    new_model = pickle.load(f)

```

predicting the output

```
test_set_pred = new_model.predict(test_set_tf)
print("So total energy consumed by the entered data will be:
{:.2f}".format(float(test_set_pred)))
```

So total energy consumed by the entered data will be: 1.54